



Faster Python Programs through Optimization

A Tutorial at EuroPython 2012

July 4, 2012

Florence, Italy

author: Dr.-Ing. Mike Müller
email: mmueller@python-academy.de
version: 3.1

Contents

1	How Fast is Fast Enough?	4
1.1	Introduction	4
1.2	Optimization Guidelines	4
2	Strategy	6
2.1	Measuring in Stones	6
2.2	Profiling CPU Usage	7
2.3	A Picture is Worth a Thousand Words	11
2.4	Going Line-by-Line	13
2.5	Profiling Memory Usage	19
2.5.1	Heapy	19
2.5.2	Pympler	25
3	Algorithms and Anti-patterns	33
3.1	String Concatenation	33
3.1.1	Exercise	33
3.2	List and Generator Comprehensions	33
3.2.1	Exercise	33
3.3	Think Global buy Local	34
4	The Right Data Structure	36
4.1	Use built-in Data Types	36
4.2	<code>list</code> vs. <code>set</code>	36
4.3	<code>list</code> vs. <code>deque</code>	40
4.4	<code>dict</code> vs. <code>defaultdict</code>	43
4.5	Big-O notation and Data Structures	46
4.6	$O(1)$ vs. $O(n)$ vs. $O(n^2)$	46
4.7	Exercises	48
5	Caching	49
5.1	Reuse before You Recalculate	49
5.2	Deterministic caching	49
5.3	Non-deterministic caching	51
5.4	Memcached	53
6	The Example	54
7	Testing Speed	55
8	Pure Python	57
8.1	Exercise	59
9	Meet Psyco, the JIT	60

9.1 Exercises	60
10 PyPy: New, Fast, Maturing	61
10.1 Exercises	61
11 NumPy for Numeric Arrays	62
11.1 Exercise	62
12 NumPyPy	63
13 Using Multiple CPUs with multiprocessing	64
14 Combination of Optimization Strategies	70
14.1 Exercise	73
15 Compilation of Tools for Speedup and Extending	74
16 End	75
16.1 Colophon	75
16.2 Links	75

1 How Fast is Fast Enough?

1.1 Introduction

Since Python is an interpreted language, some types of computations are slower in Python than in compiled languages. Depending on the application, this may or may not be a problem. This tutorial introduces several methods to speed up Python. Before starting to optimize, however the cost involved should be considered. Optimized code may need more effort to develop and maintain, leading to prolonged development time. So there is always a balance between speed of development and speed of program execution.

1.2 Optimization Guidelines

Premature optimization is the root of all evil.

C. A. R. Hoare (often misattributed to D. Knuth)

Before you start thinking about optimization make sure your program works correctly. Never optimize before the program produces the desired results.

Optimization often comes with a price: It tends to make your code less readable. Since most of the programming time for software is spent on maintenance rather than developing new code, readability and maintainability is of great importance for an effective life cycle of your program. Therefore, always think twice if it is really worth before you make your code less readable the speed gain. After all, we deliberately choose Python for its excellent readability and pay with somewhat slower programs for certain tasks.

A few general guidelines are formulated as follows:

1. Make sure your program is really too slow. Do you really need more performance? Are there any other slowdown factors such as network traffic or user input that have more impact on speed? Does it hurt if the program runs slowly?
2. Don't optimize as you go. Don't waste time before you are certain that you will need the additional speed.
3. Only realistic use cases and user experience should be considered.
4. Architecture can be essential for performance. Is it appropriate?
5. Are there any bugs that slow down the program?
6. If the program is really too slow, find the bottlenecks by profiling (use module `profile`).
7. Always check the result of optimization with all unit tests. Don't optimize with bugs.

Usually for complex programs, the most performance gain can be achieved by optimization of algorithms. Finding what big-O notation an algorithm has is very important to predict performance for large amounts of data.

The first thing you should check before making any changes in your program is external causes that may slow down your program. Likely candidates are:

- network connections
- database access
- calls to system functions

In most cases hardware is cheaper than programmer time. Always check if there is enough memory for application. Swapping memory pages to disc may slow down execution by an order of magnitude. Make sure you have plenty of free disk space and a recent and fast processor. The Python Cookbook [MART2005], also available [online](#) ¹, is a very good compilation of short and not so short solutions to

1 How Fast is Fast Enough?

specific problems. Some of the recipes, especially in the algorithm section are applicable to performance issues.

The Python in a Nutshell book ([MART2006]) contains a good summary on optimization, including profiling as well as large-scale and small-scale optimization (see pages 474 to 489). There are two chapters about optimization in [ZIAD2008]. A good resource for scientific applications in Python is [LANG2006] that also contains substantial material on optimization and extending of Python with other languages.

Some of them are exemplified in the following section.

From now on we assume you have done all the above-mentioned steps and still need more speed.

2 Strategy

2.1 Measuring in Stones

Programs will run at different speeds on different hardware. The use of benchmarks allows to measure how fast your hardware and in the case of Python how fast the used implementation is. Python has the module `test.pystone` that allows to benchmark hardware and implementation. We can use it as a standalone script:

```
python2.5 ..\Python25\Lib\test>pystone.py
Pystone(1.1) time for 50000 passes = 0.840635
This machine benchmarks at 59478.8 pystones/second
```

Python 2.6 is a little bit faster:

```
python2.6 ..\Python26\Lib\test>pystone.py
Pystone(1.1) time for 50000 passes = 0.6943
This machine benchmarks at 72015 pystones/second
```

Python 2.7 needs a bit longer for this test:

```
python2.7 ..\Python27\Lib\test>pystone.py
Pystone(1.1) time for 50000 passes = 0.807359
This machine benchmarks at 61930.3 pystones/second
```

IronPython is faster than CPython for this benchmark:

```
ipy ..\Lib\test\pystone.py
Pystone(1.1) time for 50000 passes = 0.459697
This machine benchmarks at 108767 pystones/second
```

PyPy is **significantly** faster than CPython and IronPython for this benchmark:

```
Pystone(1.1) time for 50000 passes = 0.152618
This machine benchmarks at 327616 pystones/second
```

And repeating the test several times will make it even faster:

```
Pystone(1.1) time for 50000 passes = 0.115852
This machine benchmarks at 431583 pystones/second
```

But Jython is much slower:

```
jython2.2 ..\Lib\test\pystone.py
Pystone(1.1) time for 50000 passes = 1.359
This machine benchmarks at 36791.8 pystones/second

jython2.5 ..\Lib\test\pystone.py
Pystone(1.1) time for 50000 passes = 1.37963
This machine benchmarks at 36241.6 pystones/second
```

2.2 Profiling CPU Usage

We can also use `pystone` in our programs:

```
>>> from test import pystone
>>> pystone.pystones()
(1.2585885652668103, 39727.04136987008)
```

The first value is the benchmark time in seconds and the second the pystones. We can use the `pystone` value to convert measured run times in seconds into pystones:

```
# file: pystone_converter.py

"""Convert seconds to kilo pystones."""

from test import pystone

BENCHMARK_TIME, PYSTONES = pystone.pystones()

def kpystone_from_seconds(seconds):
    """Convert seconds to kilo pystones."""
    return (seconds * PYSTONES) / 1e3

if __name__ == '__main__':

    def test():
        """Show how it works"""

        print
        print '%10s %10s' % ('seconds', 'kpystones')
        print
        for seconds in [0.1, 0.5, 1.0, 2.0, 5.0]:
            print ('%10.5f %10.5f' % (seconds, kpystone_from_seconds(seconds)))

    test()
```

We will use this function to compare our results.

2.2 Profiling CPU Usage

There are three modules in the Python standard library that allow measuring the used CPU time:

- `profile`
- `hotshot` and
- `cProfile`

Because `profile` is a pure Python implementation and `hotshot` might be removed in a future version of Python, `cProfile` is the recommended tool. It is part of the standard library for version 2.5 onwards. All three profilers are deterministic and therefore actually run the code they are profiling and measure its execution time. This has some overhead but provides reliable results in most cases. `cProfile` tries to minimize this overhead. Since Python works with the interpreter, the overhead is rather small. The other type of profiling is called statistical and uses random sampling of the effective instruction pointer. This has less overhead but is also less precise. We won't look at those techniques.

Let's write a small program whose whole purpose is to use up CPU time:

2.2 Profiling CPU Usage

```
# file profile_me.py

"""Example to be profiled.
"""

import time

def fast():
    """Wait 0.001 seconds.
    """
    time.sleep(1e-3)

def slow():
    """Wait 0.1 seconds.
    """
    time.sleep(0.1)

def use_fast():
    """Call `fast` 100 times.
    """
    for _ in xrange(100):
        fast()

def use_slow():
    """Call `slow` 100 times.
    """
    for _ in xrange(100):
        slow()

if __name__ == '__main__':
    use_fast()
    use_slow()
```

Now we import our module as well as cProfile:

```
>>> import profile_me
>>> import cProfile
```

and make an instance of Profile:

```
>>> profiler = cProfile.Profile()
```

First we call our fast function:

```
>>> profiler.runcall(profile_me.use_fast)
```

and look at the statistics cProfile provides:

2.2 Profiling CPU Usage

```
>>> profiler.print_stats()
      202 function calls in 0.195 CPU seconds
Ordered by: standard name
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
    100     0.000     0.000     0.195     0.002 profile_me.py:3(fast)
         1     0.000     0.000     0.195     0.195 profile_me.py:9(use_fast)
         1     0.000     0.000     0.000     0.000 ~:0(<method 'disable' of
         ' _lsprof.Profiler' objects>)
    100     0.194     0.002     0.194     0.002 ~:0(<time.sleep>)
```

The column headers have the following meaning:

- `ncalls` is the number of calls to this function
- `tottime` is the total time spent in this function, where calls to sub-functions are excluded from time measurement
- `percall` is `tottime` divided by `ncalls`
- `cumtime` is the cumulative time, that is the total time spent in this including the time spent in sub-functions
- `percall` is `cumtime` divided by `ncalls`
- `filename:lineno(function)` are the name of the module, the line number and the name of the function

We can see that the function `fast` is called 100 times and that it takes about 0.002 seconds per call. At first look it is surprising that `tottime` is zero. But if we look at the time the function `time.sleep` uses up, it becomes clear the `fast` spends only 0.001 seconds ($0.195 - 0.194$ seconds) and the rest of the time is burnt in `time.sleep()`.

We can do the same thing for our slow function:

```
>>> profiler = cProfile.Profile()
>>> profiler.runcall(profile_me.use_slow)
>>> profiler.print_stats()
      202 function calls in 10.058 CPU seconds
Ordered by: standard name
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
         1     0.001     0.001    10.058    10.058 profile_me.py:13(use_slow)
    100     0.001     0.000    10.058     0.101 profile_me.py:6(slow)
         1     0.000     0.000     0.000     0.000 ~:0(<method 'disable' of
         ' _lsprof.Profiler' objects>)
    100    10.057     0.101    10.057     0.101 ~:0(<time.sleep>)
```

Not surprisingly, the run times are nearly two orders of magnitude greater, because we let `sleep` use up one hundred times more time.

Another method to invoke the profiler is to use the function `run`:

```
>>> cProfile.run('profile_me.use_fast()')
      203 function calls in 0.195 CPU seconds
Ordered by: standard name
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
         1     0.000     0.000     0.195     0.195 <string>:1(<module>)
    100     0.000     0.000     0.195     0.002 profile_me.py:3(fast)
```

2.2 Profiling CPU Usage

```
1      0.000      0.000      0.195      0.195 profile_me.py:9(use_fast)
1      0.000      0.000      0.000      0.000 ~:0(<method 'disable' of
      '_lsprof.Profiler' objects>)
100    0.195      0.002      0.195      0.002 ~:0(<time.sleep>)
```

Here we supply the function to be called as a string with parenthesis, i.e. a string that can be used in an `exec` statement as opposed to the function object we supplied to the `runcall` method of our `Profile` instance.

We can also supply a file where the measured runtime data will be stored:

```
>>> cProfile.run('profile_me.use_fast()', 'fast.stats')
```

Now we can use the `pstats` module to analyze these data:

```
>>> cProfile.run('profile_me.use_fast()', 'fast.stats')
>>> import pstats
>>> stats = pstats.Stats('fast.stats')
```

We can just print out the data in the same format we saw before:

```
>>> stats.print_stats()
Wed Mar 11 16:11:39 2009      fast.stats
      203 function calls in 0.195 CPU seconds
      Random listing order was used
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
   100    0.194    0.002    0.194    0.002 ~:0(<time.sleep>)
     1     0.000    0.000    0.000    0.000 ~:0(<method 'disable' of '_lsprof.Profiler'
objects>)
   100    0.000    0.000    0.195    0.002 profile_me.py:3(fast)
     1     0.000    0.000    0.195    0.195 <string>:1(<module>)
     1     0.000    0.000    0.195    0.195 profile_me.py:9(use_fast)
```

We can also sort by different columns and restrict the number of lines printed out. Here we sort by the number of calls and want to see only the first three columns:

```
>>> stats.sort_stats('calls').print_stats(3)
Wed Mar 11 16:11:39 2009      fast.stats
      203 function calls in 0.195 CPU seconds
      Ordered by: call count
      List reduced from 5 to 3 due to restriction <3>
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
   100    0.194    0.002    0.194    0.002 ~:0(<time.sleep>)
   100    0.000    0.000    0.195    0.002 profile_me.py:3(fast)
     1     0.000    0.000    0.000    0.000 ~:0(<method 'disable' of
      '_lsprof.Profiler' objects>)
```

Or we sort by time used and show all lines:

```
>>> stats.sort_stats('time').print_stats()
Wed Mar 11 16:11:39 2009      fast.stats
      203 function calls in 0.195 CPU seconds
      Ordered by: internal time
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
```

2.3 A Picture is Worth a Thousand Words

```
100      0.194      0.002      0.194      0.002 ~:0(<time.sleep>)
100      0.000      0.000      0.195      0.002 profile_me.py:3(fast)
  1      0.000      0.000      0.195      0.195 profile_me.py:9(use_fast)
  1      0.000      0.000      0.195      0.195 <string>:1(<module>)
  1      0.000      0.000      0.000      0.000 ~:0(<method 'disable' of
      '_lsprof.Profiler' objects>)
```

We can also get information about which function is called by a certain function:

```
>>> stats.print_callers('fast')
Ordered by: internal time
List reduced from 5 to 2 due to restriction <'fast'>
Function                                was called by...
profile_me.py:3(fast)                    profile_me.py:9(use_fast)
((100, 100, 0.00040628818660897912, 0 .19478914258667296))    0.195
profile_me.py:9(use_fast)                  <string>:1(<module>)
((1, 1, 0.00026121123840443721, 0.1950503538250774))    0.195
```

We can also find out what functions are called:

```
>>> stats.print_callees('use_fast')
Ordered by: internal time
List reduced from 5 to 1 due to restriction <'use_fast'>
Function                                called...
profile_me.py:9(use_fast)                profile_me.py:3(fast)
((100, 100, 0.00040628818660897912, 0.19478914258667296))    0.195
```

There are more interesting attributes such as the number of calls:

```
>>> stats.total_calls
203
```

2.3 A Picture is Worth a Thousand Words


Doing the statistics with tables is worthwhile and interesting. But there is another way to look at the profiling results: making graphs. A very nice tool for this is [RunSnakeRun](#)². It is written in Python itself and uses wxPython and SquareMap.

The usage is very simple. After installing RunSnakeRun just type:


```
runsnake slow.stats
```

at the command line and you will get nice interactive graphs that should look like this for our slow example:

2.3 A Picture is Worth a Thousand Words



Our fast example is not really fast and the graphical view shows a very similar picture

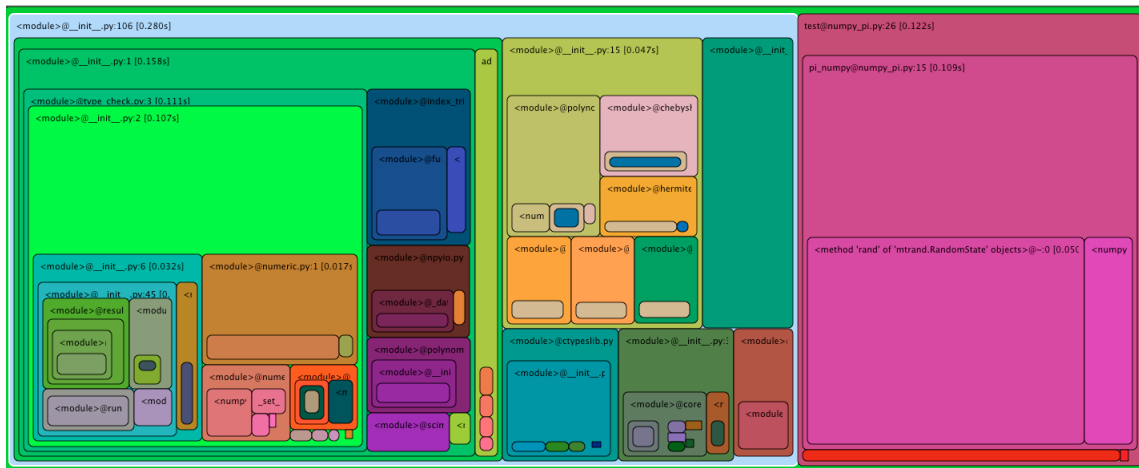


Later in the course, we will use a simple algorithm to calculate pi (the one from the circle, see chapter "The Example" for more details and the code). This gives a more interesting picture:

2.4 Going Line-by-Line



We also will use a NumPy version for this algorithm. Even though our code has about the same number of lines, the graph becomes much more complex because we use NumPy functions:



2.4 Going Line-by-Line

With `cProfile` the finest resolution we get is the function call. But there is `line_profiler` by Robert Kern that allows line-by-line profiling. `line_profiler` comes in bundle with `kernprof` that adds some features to `cProfile`. Installation is simple:

```
pip install line_profiler
```

We can use `kernprof` from the command line, which turns just uses `cProfile`. The option `-v` shows the statics right away:

```
$ kernprof.py -v profile_me.py
Wrote profile results to profile_me.py.prof
      406 function calls in 10.204 seconds
```

Ordered by: standard name

2.4 Going Line-by-Line

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	10.204	10.204	<string>:1(<module>)
100	0.001	0.000	10.081	0.101	profile_me.py:15(slow)
1	0.001	0.001	0.121	0.121	profile_me.py:21(use_fast)
1	0.001	0.001	10.082	10.082	profile_me.py:28(use_slow)
1	0.001	0.001	10.204	10.204	profile_me.py:4(<module>)
100	0.001	0.000	0.120	0.001	profile_me.py:9(fast)
1	0.000	0.000	10.204	10.204	{execfile}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}
200	10.199	0.051	10.199	0.051	{time.sleep}

We add the decorator `profile` to the function we would like to profile:

```
# file profile_me_use_line_profiler.py

"""Example to be profiled.
"""

import time

def fast():
    """Wait 0.001 seconds.
    """
    time.sleep(1e-3)

def slow():
    """Wait 0.1 seconds.
    """
    time.sleep(0.1)

@profile
def use_fast():
    """Call `fast` 100 times.
    """
    for _ in xrange(100):
        fast()

@profile
def use_slow():
    """Call `slow` 100 times.
    """
    for _ in xrange(100):
        slow()

if __name__ == '__main__':
    use_fast()
    use_slow()
```

Now we can use the option `-l` to turn on `line_profiler`:

```
$ kernprof.py -l -v profile_me_use_line_profiler.py
Wrote profile results to profile_me_use_line_profiler.py.lprof
```

2.4 Going Line-by-Line

Timer unit: 1e-06 s

File: profile_me_use_line_profiler.py

Function: use_fast at line 20

Total time: 0.120634 s

Line #	Hits	Time	Per Hit	% Time	Line Contents
20					@profile
21					def use_fast():
22					"""Call `fast` 100 times.
23					"""
24	101	732	7.2	0.6	for _ in xrange(100):
25	100	119902	1199.0	99.4	fast()

File: profile_me_use_line_profiler.py

Function: use_slow at line 27

Total time: 10.086 s

Line #	Hits	Time	Per Hit	% Time	Line Contents
27					@profile
28					def use_slow():
29					"""Call `slow` 100 times.
30					"""
31	101	1147	11.4	0.0	for _ in xrange(100):
32	100	10084845	100848.4	100.0	slow()

This shows us how much time each line used. Our test functions are very short. Let's create a small function that accumulates the sums of all elements in a list.

```
"""Simple test function for line_profiler.
"""

@profile
def accumulate(iterable):
    """Accumulate the intermediate steps in summing all elements.

    The result is a list with the lenght of `iterable`.
    The last elments is the sum of all elements of `ieterable`
    >>>accumulate(range(5))
    [0, 1, 3, 6, 10]
    accumulate(range(10))
    [0, 1, 3, 6, 10, 15, 21, 28, 36, 45]
    """
    acm = [iterable[0]]
    for elem in iterable[1:]:
        old_value = acm[-1]
        new_value = old_value + elem
        acm.append(new_value)
    return acm

if __name__ == '__main__':
```

2.4 Going Line-by-Line

```
accumulate(range(10))
accumulate(range(100))
```

Let's look at the output:

```
$ kernprof.py -l -v accumulate.py
Wrote profile results to accumulate.py.lprof
Timer unit: 1e-06 s

File: accumulate.py
Function: accumulate at line 3
Total time: 0.000425 s

Line #      Hits          Time Per Hit   % Time  Line Contents
=====
   3                                     @profile
   4                                     def accumulate(iterable):
   5                                     """Accumulate the intermediate steps in summing all elements.
   6
   7                                     The result is a list with the lenght of `iterable`.
   8                                     The last elements is the sum of all elements of `ieterable`
   9                                     >>>accumulate(range(5))
  10                                     [0, 1, 3, 6, 10]
  11                                     accumulate(range(10))
  12                                     [0, 1, 3, 6, 10, 15, 21, 28, 36, 45]
  13                                     """
  14                                     acm = [iterable[0]]
  15                                     for elem in iterable[1:]:
  16                                         old_value = acm[-1]
  17                                         new_value = old_value + elem
  18                                         acm.append(new_value)
  19                                     return acm
```

The algorithm could be written more concise. In fact the three lines inside the loop could be one. But we would like to see how much each operation takes and therefore spread things over several lines.

Another example looks at some simple mathematical calculations:

```
"""Simple test function for line_profiler doing some math.
"""

import math

@profile
def calc(number, loops=1000):
    """Do some math calculations.
    """
    sqrt = math.sqrt
    for x in xrange(loops):
        x = number + 10
        x = number * 10
        x = number ** 10
        x = pow(x, 10)
        x = math.sqrt(number)
        x = sqrt(number)
        math.sqrt
        sqrt

if __name__ == '__main__':
    calc(100, int(1e5))
```

The output shows which operation takes the most time:

2.4 Going Line-by-Line

```
$ kernprof.py -l -v calc.py
Wrote profile results to calc.py.lprof
Timer unit: 1e-06 s
```

```
File: calc.py
Function: calc at line 7
Total time: 1.33158 s
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
=====					
7					@profile
8					def calc(number, loops=1000):
9					"""Do some math calculations.
10					"""
11	1	4	4.0	0.0	sqrt = math.sqrt
12	100001	77315	0.8	5.8	for x in xrange(loops):
13	100000	87124	0.9	6.5	x = number + 10
14	100000	84518	0.8	6.3	x = number * 10
15	100000	330587	3.3	24.8	x = number ** 10
16	100000	378584	3.8	28.4	x = pow(x, 10)
17	100000	109849	1.1	8.2	x = math.sqrt(number)
18	100000	93211	0.9	7.0	x = sqrt(number)
19	100000	88768	0.9	6.7	math.sqrt
20	100000	81624	0.8	6.1	sqrt

The function `pow` takes by far the most time, whereas `sqrt` from the `math` module is fast. Note that seem to be no difference between `math.sqrt` and `sqrt`, which is just a local reference. Let's look at this in a further example:

```
"""Testing access to local name and name referenced on another module.
"""

import math

# If there is no decorator `profile`, make one that just calls the function,
# i.e. does nothing.
# This allows to call `kernprof` with and without the option `-l` without
# commenting or un-commenting `@profile` all the time.
# You can add this to the builtins to make it available in the whole program.
try:
    @profile
    def dummy():
        """Needs to be here to avoid a syntax error.
        """
        pass
except NameError:
    def profile(func):
        """Will act as the decorator `profile` if it was already found.
        """
        def mock(*args, **kwargs):
            """Just call the function. No actual decoration effect.
            """
            return func(*args, **kwargs)
        return mock
```

2.4 Going Line-by-Line

```
def local_ref(counter):
    """Access local name.
    """
    # make it local
    sqrt = math.sqrt
    for _ in xrange(counter):
        sqrt

def module_ref(counter):
    """Access name as attribute of another module.
    """
    for _ in xrange(counter):
        math.sqrt

@profile
def test(counter):
    """Call both functions.
    """
    local_ref(counter)
    module_ref(counter)

if __name__ == '__main__':
    test(int(1e8))
```

There are two functions to be line-traced. `local_ref` gets a local reference to `math.sqrt` and `module_ref` calls `math.sqrt` as it is.

We run this with the option `-v`, and we get:

```
$ kernprof.py -v local_ref.py
Wrote profile results to local_ref.py.prof
  9 function calls in 14.847 seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
      1      0.000      0.000    14.847    14.847  <string>:1(<module>)
      1      0.000      0.000      0.000      0.000  local_ref.py:18(profile)
      1      0.001      0.001    14.846    14.846  local_ref.py:2(<module>)
      1      0.000      0.000    14.845    14.845  local_ref.py:21(mock)
      1      4.752      4.752      4.752      4.752  local_ref.py:28(local_ref)
      1     10.093     10.093     10.093     10.093  local_ref.py:37(module_ref)
      1      0.000      0.000    14.845    14.845  local_ref.py:44(test)
      1      0.001      0.001    14.847    14.847  {execfile}
      1      0.000      0.000      0.000      0.000  {method 'disable' of
'_lsprof.Profiler' objects}
```

This shows that `local_ref` is more than twice as fast as `module_ref` because it avoids many lookups on the module `math`.

Now we run it with the options `-v - 1`:

2.5 Profiling Memory Usage

```
$ kernprof.py -v -l local_ref.py
Wrote profile results to local_ref.py.lprof
Timer unit: 1e-06 s

File: local_ref.py
Function: dummy at line 12
Total time: 0 s

Line #      Hits          Time Per Hit   % Time  Line Contents
=====
    12                               @profile
    13                               def dummy():
    14                                   """Needs to be here to
avoid a syntax error.
    15                                   """
    16                                   pass

File: local_ref.py
Function: test at line 44
Total time: 125.934 s

Line #      Hits          Time Per Hit   % Time  Line Contents
=====
    44                               @profile
    45                               def test(counter):
    46                                   """Call both functions.
    47                                   """
    48                               1      58162627 58162627.0    46.2      local_ref(counter)
    49                               1      67771433 67771433.0    53.8      module_ref(counter)
```

This takes much longer. The differences in run times are largely gone. After correspondence with Robert Kern, the author of `line_profiler`, it turns out that the substantial overhead the line tracing adds causes a distortion of measuring results. Conclusion: Use `line_profiler` for expensive atomic calls such as to a function in an extension module like NumPy.

2.5 Profiling Memory Usage

Current computers have lots of RAM, still it can be a problem if an application uses more RAM than is physically available, leading to swapping and a large performance penalty. In particular, long running applications tend to use up more RAM over time. Although Python does automatic memory management, there are cases where memory is not released because there are still references to objects that are no longer needed. We can call the garbage collector manually, but this does not always produce the desired effects.

2.5.1 Heapy

The [Guppy_PE framework](#) ³ provides the tool `heapy` that is very useful for inspecting Python memory usage. It is not the easiest tool to work with but still provides valuable insights in how memory is used by Python objects.

Unfortunately, `heapy` works only up to Python 2.6. The most recent update seems to be from 2009. There was a change in the Python C API and `_PyLong_AsScaledDouble` was removed.

We look at some features at the interactive prompt. First we import `hpy` and call it:

2.5 Profiling Memory Usage

```
>>> from guppy import hpy
>>> hp = hpy()
```

Now we can look at the heap:

```
>>> hp.heap()
Partition of a set of 55690 objects. Total size = 3848216 bytes.
  Index  Count    %      Size  % Cumulative  % Kind (class / dict of class)
    0    27680   50   1522412   40   1522412   40 str
    1      150    0    666120   17   2188532   57 dict of module
    2   10459   19    474800   12   2663332   69 tuple
    3    2913    5    186432    5   2849764   74 types.CodeType
    4    2814    5    168840    4   3018604   78 function
    5     368    1    167488    4   3186092   83 dict (no owner)
    6     345    1    151596    4   3337688   87 dict of class
    7     145    0     90956    2   3428644   89 dict of type
    8     192    0     82156    2   3510800   91 type
    9    6310   11     75720    2   3586520   93 int
<140 more rows. Type e.g. '_.more' to view.>
```

There are 150 types of objects in our fresh interactive session. 40 % of the memory is taken up by strings and 17 % by module name space dictionaries.

We create a new object, a list with one million intergers:

```
>>> big_list = range(int(1e6))
```

and look at our heap again:

```
>>> hp.heap()
Partition of a set of 1055602 objects. Total size = 19924092 bytes.
  Index  Count    %      Size  % Cumulative  % Kind (class / dict of class)
    0 1006210   95 12074520   61 12074520   61 int
    1     208    0  4080320   20 16154840   81 list
    2   27685    3  1522628    8 17677468   89 str
    3     150    0    666120    3 18343588   92 dict of module
    4   10458    1    474768    2 18818356   94 tuple
    5    2913    0    186432    1 19004788   95 types.CodeType
    6    2813    0    168780    1 19173568   96 function
    7     374    0    168328    1 19341896   97 dict (no owner)
    8     345    0    151596    1 19493492   98 dict of class
    9     145    0     90956    0 19584448   98 dict of type
<140 more rows. Type e.g. '_.more' to view.>
```

Now integers, of which we have one million in our list, take up 61 % of the memory followed by lists that use up 20 %. Strings are down to 8%. We delete our list:

```
>>> del big_list
```

and we are (nearly) back to our initial state:

2.5 Profiling Memory Usage

```
>>> hp.heap()
Partition of a set of 55700 objects. Total size = 3861984 bytes.
  Index  Count    %      Size  % Cumulative  % Kind (class / dict of class)
    0    27685  50   1522632  39   1522632  39 str
    1      150   0    666120  17   2188752  57 dict of module
    2   10458  19    474768  12   2663520  69 tuple
    3    2913   5    186432   5   2849952  74 types.CodeType
    4    2813   5    168780   4   3018732  78 function
    5     374   1    168328   4   3187060  83 dict (no owner)
    6     345   1    151596   4   3338656  86 dict of class
    7     145   0     90956   2   3429612  89 dict of type
    8     192   0     82156   2   3511768  91 type
    9    6309  11     75708   2   3587476  93 int
<140 more rows. Type e.g. '_.more' to view.>
```

We can tell `hp` to count only newly added objects with:

```
>>> hp.setref()
```

There are still a few objects, but much fewer than before:

```
>>> hp.heap()
Partition of a set of 93 objects. Total size = 8768 bytes.
  Index  Count    %      Size  % Cumulative  % Kind (class / dict of class)
    0      8   9     4024  46     4024  46 types.FrameType
    1      7   8      980  11     5004  57 dict of type
    2     16  17      704   8     5708  65 __builtin__.weakref
    3     20  22      700   8     6408  73 tuple
    4      4   4      560   6     6968  79 dict (no owner)
    5      4   4      560   6     7528  86 dict of guppy.etc.Glue.Interface
    6      9  10      320   4     7848  90 str
    7      7   8      280   3     8128  93 __builtin__.wrapper_descriptor
    8      1   1      140   2     8268  94 dict of guppy.etc.Glue.Owner
    9      4   4      128   1     8396  96 guppy.etc.Glue.Interface
<5 more rows. Type e.g. '_.more' to view.>
```

Now we can create our big list again:

```
>>> big_list = range(int(1e6))
```

The list and the integers in it take up 99 % (74 + 25) of the memory now:

```
>>> hp.heap()
Partition of a set of 1000742 objects. Total size = 16120680 bytes.
  Index  Count    %      Size  % Cumulative  % Kind (class / dict of class)
    0  999908 100  11998896  74  11998896  74 int
    1      3   0    4066700  25  16065596 100 list
    2     750   0     46532   0  16112128 100 str
    3      8   0      4012   0  16116140 100 types.FrameType
    4      7   0       980   0  16117120 100 dict of type
    5     22   0       776   0  16117896 100 tuple
    6     16   0       704   0  16118600 100 __builtin__.weakref
```

2.5 Profiling Memory Usage

```
7      4      0      560      0 16119160 100 dict (no owner)
8      4      0      560      0 16119720 100 dict of guppy.etc.Glue.Interface
9      7      0      280      0 16120000 100 __builtin__.wrapper_descriptor
<8 more rows. Type e.g. '_.more' to view.>
```

Even we have an error of 1 % in our example, it is good enough to find out how memory changes when we do certain things.

If we use `setref` several times in a row, we get slightly different results:

```
>>> h.size
16120804
>>> hp.setref()
>>> h.size
16120804
>>> hp.heap().size
5620
>>> big_list = range(int(1e6))
>>> hp.heap().size
16067724
>>> hp.setref()
>>> hp.heap().size
4824
>>> big_list = range(int(1e6))
>>> hp.heap().size
16066788
>>> hp.setref()
>>> hp.heap().size
4768
```

There is much more information in the heap. Let's look at some of them:

```
>>> h = hp.heap()
```

We can use the index to extract single lines:

```
>>> h[0]
Partition of a set of 999910 objects. Total size = 11998920 bytes.
Index  Count    %      Size    % Cumulative % Kind (class / dict of class)
0 999910 100 11998920 100 11998920 100 int
```

We can order everything by type:

```
>>> h.bytype
Partition of a set of 1000746 objects. Total size = 16120804 bytes.
Index  Count    %      Size    % Cumulative % Type
0 999910 100 11998920 74 11998920 74 int
1      3      0  4066700 25 16065620 100 list
2    750      0   46536      0 16112156 100 str
3      8      0    4028      0 16116184 100 types.FrameType
4     17      0    2380      0 16118564 100 dict
5     24      0     856      0 16119420 100 tuple
6     16      0     704      0 16120124 100 __builtin__.weakref
```

2.5 Profiling Memory Usage

```
7      7      0      280      0 16120404 100 __builtin__.wrapper_descriptor
8      4      0      128      0 16120532 100 guppy.etc.Glue.Interface
9      3      0      120      0 16120652 100 types.MethodType
<3 more rows. Type e.g. '_.more' to view.>
```

Since there are only three more lines to display, we use the method `more` to see all of `h` content:

```
>>> _.more
Index  Count  %      Size  % Cumulative  % Type
10      2      0       72      0 16120724 100 types.InstanceType
11      1      0       64      0 16120788 100 types.CodeType
12      1      0       16      0 16120804 100 long
```

We can also order by referrers:

```
>>> h.byrcs
Partition of a set of 1000746 objects. Total size = 16120804 bytes.
Index  Count  %      Size  % Cumulative  % Referrers by Kind (class / dict of class)
0 1000648 100 12045316 75 12045316 75 list
1      3      0 4063336 25 16108652 100 dict of module
2     27      0   4708      0 16113360 100 <Nothing>
3      6      0   3472      0 16116832 100 tuple
4     21      0   1456      0 16118288 100 type
5      4      0    560      0 16118848 100 guppy.etc.Glue.Interface
6      3      0    420      0 16119268 100 dict of guppy.etc.Glue.Owner
7      8      0    352      0 16119620 100 guppy.heapy.heapyc.HeapView
8      7      0    280      0 16119900 100 dict of type
9      7      0    256      0 16120156 100 dict (no owner), dict of guppy.etc.Glue.Interface
<9 more rows. Type e.g. '_.more' to view.>
```

Let's look at some examples for how we can use `hpy`. First we write a decorator that tells us how much memory the result of a function uses:

```
# file: memory_size_hpy.py

"""Measure the size of used memory with a decorator.
"""

import functools #1

from guppy import hpy #2

memory = {} #3

def measure_memory(function): #4
    """Decorator to measure memory size.
    """

    @functools.wraps(function) #5
    def _measure_memory(*args, **kwargs): #6
        """This replaces the function that is to be measured.
        """

        measurer = hpy() #7
        measurer.setref() #8
```

2.5 Profiling Memory Usage

```
    initial_memory = measurer.heap().size #9
    try:
        res = function(*args, **kwargs) #10
        return res
    finally:
        memory[function.__name__] = (measurer.heap().size - #11
                                     initial_memory)
    return _measure_memory #12

if __name__ == '__main__':

    @measure_memory #13
    def make_big(number):
        """Example function that makes a large list.
        """
        return range(number) #14

    make_big(int(1e6)) #15
    print 'used memory', memory #16
```

First we import `functools` (#1) that will help us to write a nice decorator. Then we import `hpy` (#2) and define a global dictionary (#3) that will hold all values for memory. We define a function that takes a function as argument (#4) and another function inside it that takes a variable number of positional and keyword arguments (#6). This is a typical setup of a decorator that takes no arguments (with arguments we would need a third level). We also decorate this function with `@functools.wraps` (#5) to preserve the docstring and the name of the original function after it is decorated.

Now we call `hpy` (#7) and set the measured memory back (#8). We measure our initially used memory (#9) and call the function with the supplied arguments (#10). We always want to have the size of memory after the call (#11). Finally, we return our internally defined function. Note that we store the result of the called function in `res`. This is necessary to get the memory that is used by the object the function returns. We return our newly created function (#12).

We decorate our function (#13) that just returns a list of size `number` (#14). After we call the function (#15), we can print the used memory (#16).

When we suspect that a function leaks memory, we can use `guppy` to measure the memory growth after a function returned:

```
# file memory._growth_hpy.py

"""Measure the memory growth during a function call.
"""

from guppy import hpy #1

def check_memory_growth(function, *args, **kwargs): #2
    """Measure the memory usage of `function`.
    """
    measurer = hpy() #3
    measurer.setref() #4
    initial_memory = measurer.heap().size #5
    function(*args, **kwargs) #6
```


2.5 Profiling Memory Usage

```
    return measurer.heap().size - initial_memory #7

if __name__ == '__main__':

    def test():
        """Do some tests with different memory usage patterns.
        """

        def make_big(number): #8
            """Function without side effects.

            It cleans up all used memory after it returns.
            """
            return range(number)

        data = [] #9

        def grow(number):
            """Function with side effects on global list.
            """
            for x in xrange(number):
                data.append(x) #10
            size = int(1e6)
            print 'memory make_big:', check_memory_growth(make_big,
                                                            size) #11
            print 'memory grow:', check_memory_growth(grow, size) #12

        test()
```

After importing `hpy` (#1) we define a helper function that takes the function to be measured, and positional and keyword arguments that will be handed to this function (#2). Now we call `hpy` (3) and set the measured memory back (#4). We measure our initially used memory (#5) and call the function with the supplied arguments (#6). Finally, we return difference in memory size before and after the function call (#7).

We define a function that just returns a list (#8) and thus does not increase memory size after it is finished. The size of the returned list is not measured.

We use a global list as data storage (#9) and define a second function that appends elements to this list (#10). Finally, we call our helper function with both functions as arguments (#11 and #12).

2.5.2 Pympler

[Pympler](#)⁴ it is a merge of the formerly independent projects `asizeof`, `heapmonitor`, and `muppy`. It works for Python version 2.4 through 3.2. We can use it very similarly to `heapy`.

Let's start a new interpreter and make an instance of `pympler.tracker.SummaryTracker`:

```
>>> from pympler import tracker
>>> mem_tracker = tracker.SummaryTracker()
```

We need to call `print_diff()` several times to get to the baseline:

```
>>> mem_tracker.print_diff()
types | # objects | total size
```

2.5 Profiling Memory Usage

```
===== | ===== | =====
list | 1353 | 138.02 KB
str | 1345 | 75.99 KB
int | 149 | 3.49 KB
dict | 2 | 2.05 KB
wrapper_descriptor | 8 | 640 B
weakref | 3 | 264 B
member_descriptor | 2 | 144 B
getset_descriptor | 2 | 144 B
function (store_info) | 1 | 120 B
cell | 2 | 112 B
instancemethod | -1 | -80 B
tuple | -1 | -216 B
>>> mem_tracker.print_diff()
types | # objects | total size
===== | ===== | =====
str | 2 | 97 B
list | 1 | 96 B
>>> mem_tracker.print_diff()
types | # objects | total size
===== | ===== | =====
```

Now we create our big list and look at the memory again:

```
>>> big_list = range(int(1e6))
>>> mem_tracker.print_diff()
types | # objects | total size
===== | ===== | =====
int | 999861 | 22.89 MB
list | 1 | 7.63 MB
```

Let's look at some examples for how we can use `pymppler`. First we write a decorator that tells us how much memory the result of a function uses:

```
# file: memory_size_pymppler.py

"""Measure the size of used memory with a decorator.
"""

import functools #1
from pymppler import tracker #2
memory = {} #3

def measure_memory(function): #4
    """Decorator to measure memory size.
    """

    @functools.wraps(function) #5
    def _measure_memory(*args, **kwargs): #6
        """This replaces the function that is to be measured.
```

2.5 Profiling Memory Usage

```
    """
    measurer = tracker.SummaryTracker() #7
    for _ in xrange(5): #8
        measurer.diff() #9
    try:
        res = function(*args, **kwargs) #10
        return res
    finally: #11
        memory[function.__name__] = (measurer.diff())
    return _measure_memory #12

if __name__ == '__main__':

    @measure_memory #13
    def make_big(number):
        """Example function that makes a large list.
        """
        return range(number) #14

    make_big(int(1e6)) #15
    print 'used memory', memory #16
```

First we import `functools` (#1) that will help us to write a nice decorator. Then we import `pymppler.tracker` (#2) and define a global dictionary (#3) that will hold all values for memory. We define a function that takes a function as argument (#4) and another function inside it that takes a variable number of positional and keyword arguments (#6). This is a typical setup of a decorator that takes no arguments (with arguments we would need a third level). We also decorate this function with `@functools.wraps` (#5) to preserve the docstring and the name of the original function after it is decorated.

Now we make an instance of our tracker (#7). We use a loop (#8) and call `tracker.diff()` several times (#9). Then we call the function with the supplied arguments (#10). We always want to have the size of memory after the call (#11). Finally, we return our internally defined function. Note that we store the result of the called function in `res`. This is necessary to get the memory that is used by the object the function returns. We return our newly created function (#12).

We decorate our function (#13) that just returns a list of size `number` (#14). After we call the function (#15), we can print the used memory (#16).

When we suspect that a function leaks memory, we can use `pymppler` to measure the memory growth after a function returned:

```
# file memory_growth_pymppler.py

"""Measure the memory growth during a function call.
"""

from pymppler import tracker #1

def check_memory_growth(function, *args, **kwargs): #2
    """Measure the memory usage of `function`.
    """
    measurer = tracker.SummaryTracker() #3
```

2.5 Profiling Memory Usage

```
for _ in xrange(5):                                #4
    measurer.diff()                                #5
function(*args, **kwargs)                          #6
return measurer.diff()                             #7

if __name__ == '__main__':

    def test():
        """Do some tests with different memory usage patterns.
        """

        def make_big(number):                       #8
            """Function without side effects.

            It cleans up all used memory after it returns.
            """
            return range(number)

        data = []                                    #9

        def grow(number):
            """Function with side effects on global list.
            """
            for x in xrange(number):
                data.append(x)                        #10
            size = int(1e6)
            print 'memory make_big:', check_memory_growth(make_big,
                                                            size)      #11
            print 'memory grow:', check_memory_growth(grow, size)  #12

        test()
```

After importing `pympler.tracker` (#1) we define a helper function that takes the function to be measured, and positional and keyword arguments that will be handed to this function (#2). We make an instance of `tracker.SummaryTracker` (3) and use a loop (#4) to call `measurer.diff()` several times and in this way set the baseline of memory usage (#5). We call the function with the supplied arguments (#6). Finally, we return difference in memory size before and after the function call (#7).

We define a function that just returns a list (#8) and thus does not increase memory size after it is finished. The size of the returned list is not measured.

We use a global list as data storage (#9) and define a second function that appends elements to this list (#10). Finally, we call our helper function with both functions as arguments (#11 and #12).

Pympler offers more tools. Let's look at the possibilities to measure the memory size of a given object. We would like to measure the memory size of a list as we append elements. We write a function that takes the length of the list and a function that is to be used to measure the memory size of an object:

```
# file: pympler_list_growth.py

"""Measure the size of a list as it grows.
"""

from pympler.asizeof import asizeof, fltsize
```

2.5 Profiling Memory Usage

```
def list_mem(lenght, size_func=flatsize):
    """Measure incremental memory increase of a growing list.
    """
    my_list= []
    mem = [size_func(my_list)]
    for elem in xrange(lenght):
        my_list.append(elem)
        mem.append(size_func(my_list))
    return mem
```


Now we use this function with three different functions: `pympler.asizeof.flatsize`, `pympler.asizeof.asizeof` and `sys.getsizeof`:

```
if __name__ == '__main__':
    SIZE = 1000
    SHOW = 20
    import sys
    for func in [flatsize, asizeof, sys.getsizeof]:
        mem = list_mem(SIZE, size_func=func)
        try:
            from matplotlib import pylab
            pylab.plot(mem)
            pylab.show()
        except ImportError:
            print 'matplotlib seems not be installed. Skipping the plot.'
            if SIZE > SHOW:
                limit = SHOW / 2
                print mem[:limit], '... skipping %d elements ...' % (SIZE - SHOW),
                print mem[-limit:]
            else:
                print mem
```

The code just calls our function and supplies one of the functions to measure memory size as an argument. If matplotlib is installed, it draws a graph for each call. Let's look at the resulting graphs.


Using `pympler.asizeof.flatsize` we get this kind of step diagram:

2.5 Profiling Memory Usage



We can see nicely how the list grows discontinuously. Python allocates more memory than it actually needs to append the next element. This way it can append several elements before it needs to increase its size again. These steps get bigger the bigger the list is.


Using `pympler.asizeof.asizeof` we get a different looking graph:



2.5 Profiling Memory Usage

This function also measures the size of all referenced objects. In our case all the integer that are stored in the list. Therefore, there is an continuous increase in memory size between the steps case by the list allocation.

For this simple case `sys.getsizeof` produces the sam result as `pympler.asizeof.flatsize`:



For more complex cases `pympler.asizeof.flatsize` might give different results.

We can also measure the number of allocation steps it takes when a list grows one element at a time:

```
# file: list_alloc_steps.py

"""Measure the number of memory allocation steps for a list.
"""

import sys

from pympler.asizeof import flatsize

def list_steps(lenght, size_func=sys.getsizeof):
    """Measure the number of memory alloaction steps for a list.
    """
    my_list = []
    steps = 0
    int_size = size_func(int())
    old_size = size_func(my_list)
    for elem in xrange(lenght):
        my_list.append(elem)
        new_size = sys.getsizeof(my_list)
        if new_size - old_size > int_size:
            steps += 1
        old_size = new_size
```

2.5 Profiling Memory Usage

```
    return steps

if __name__ == '__main__':
    print 'Using sys.getsizeof:'
    for size in [10, 100, 1000, 10000, int(1e5), int(1e6), int(1e7)]:
        print '%10d: %3d' % (size, list_steps(size))
    print 'Using pympler.asizeof.flatsize:'
    for size in [10, 100, 1000, 10000, int(1e5), int(1e6), int(1e7)]:
        print '%10d: %3d' % (size, list_steps(size, flatsize))
```

The results are the same for `sys.getsizeof` and `pympler.asizeof.flatsize`:

```
Using sys.getsizeof:
    10:    3
   100:   10
  1000:   27
 10000:   46
100000:   65
1000000:  85
10000000: 104
Using pympler.asizeof.flatsize:
    10:    3
   100:   10
  1000:   27
 10000:   46
100000:   65
1000000:  85
10000000: 104
```


3 Algorithms and Anti-patterns

3.1 String Concatenation

Strings in Python are immutable. So if you want to modify a string, you have to actually create a new one and use parts of the old one:

```
>>> s = 'old text'
>>> 'new' + s[-5:]
'new text'
```

This means that new memory has to be allocated for the string. This is no problem for a few hundred or thousand strings, but if you have to deal with millions of strings, memory allocation time may be considerably longer. The solution in Python is to use a list to hold the sub strings and join them with `''.join()` string method.

3.1.1 Exercise

Write a test program that constructs a very long string (containing up to one million characters). Use the idiom `s += 'text'` and the idiom `text_list.append('text')` plus `''.join(text_list)` in a function for each. Compare the two approaches in terms of execution speed.

Hint: You can use `timeit.default_timer()` to get the time since the last call to this function. Alternatively, you can use the module `timeit` or the function `measureRunTime` which can be found in the `examples` directory in the subdirectory `modules`.

3.2 List and Generator Comprehensions

Python offers list comprehension as a short and very readable way to construct a list.

```
>>> l = [x * x for x in xrange(10)]
>>> l
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
```

is a short form for:

```
>>> for x in xrange(10):
...     l.append(x * x)
...
>>> l
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
```

If you are not interested in the list itself but rather some values computed from the whole list, you can use generator comprehension and avoid the list all together.

```
>>> sum(x * x for x in xrange(10))
285
```

3.2.1 Exercise

Write a test program that calculates the sum of all squares of the numbers form zero to one million. Use the idiom `l.append` and list comprehension as well as generator comprehension. Try it with `range` and `xrange`. Use different numbers, e.g. smaller and bigger than one million.

3.3 Think Global buy Local

Hint: You can use `timeit.default_timer()` to get the time since the last call to this function. Alternatively, you can use the module `timeit` or the function `measureRunTime` which you can find in the `examples` directory in the subdirectory `modules`.

3.3 Think Global buy Local

A greta deal of things in Python are dynamic. This includes the lookup of variables. It follows the famous LGB local-global-built-in rule. If a variable name is not found in the local scope, Python looks for it in global and then in the built-in name space before raising an `NameError` when nothing was found.

Since every name space is a dictionary, it involves more look ups the more name spaces have to be searched. Therefore, local variables are faster than global variables. Let's look at an example:

```
# file: local_global.py

"""Local vs. built-in.
"""

GLOBAL = 1

def repeat(counter):
    """Using the GLOBAL value directly.
    """
    for count in xrange(counter):
        GLOBAL

def repeat_local(counter):
    """Making GLOBAL a local variable.
    """
    local = GLOBAL
    for count in xrange(counter):
        local

def test(counter):
    """Call both functions.
    """
    repeat(counter)
    repeat_local(counter)

if __name__ == '__main__':

    def do_profile():
        """Check the run times.
        """
        import cProfile
        profiler = cProfile.Profile()
        profiler.run('test(int(1e8))')
        profiler.print_stats()

    do_profile()
```

3.3 Think Global buy Local

By running this code, we will see that the version that accesses the `GLOBAL` directly is about 25% slower than the version with the local variable.

The difference becomes larger when we move more outward and make a built-in name a local one:

```
"""Local vs. built-in.
"""

def repeat(counter):
    """Using the built-in `True` in a loop.
    """
    for count in xrange(counter):
        True

def repeat_local(counter):
    """Making `True` a local variable.
    """
    true = True
    for count in xrange(counter):
        true

def test(counter):
    """Call both functions.
    """
    repeat(counter)
    repeat_local(counter)

if __name__ == '__main__':

    def do_profile():
        """Check the run times.
        """
        import cProfile
        profiler = cProfile.Profile()
        profiler.run('test(int(1e8))')
        profiler.print_stats()

    do_profile()
```

In this example it saves about 40% of the run. So, if you have large loops and you access globals or built-ins frequently, making them local might be quite useful.

4 The Right Data Structure

4.1 Use built-in Data Types

It is always a good idea to use Python built-in data structures. They are not only most often more elegant and robust than self-made data structures, but also faster in nearly all cases. They are well tested, often partially implemented in C and optimized through long time usage by many of talented programmers.

There are essential differences among built-in data types in terms of performance depending on the task.

4.2 list vs. set

If you need to search in items, dictionaries and sets are mostly preferable to lists.

```
>>> 9 in range(10)
True
>>> 9 in set(range(10))
True
```

Let's make a performance test. We define a function that searches in a list:

```
>>> import timeit
>>> def search_list(n):
...     my_list = range(n)
...     start = timeit.default_timer()
...     n in my_list
...     return timeit.default_timer() - start
... 
```

and one that searches in a set:

```
>>> def search_set(n):
...     my_set = set(range(n))
...     start = timeit.default_timer()
...     n in my_set
...     return timeit.default_timer() - start
... 
```

We define a function that compares both run time:

```
>>> def compare(n):
...     print 'ratio:', search_list(n) / search_set(n)
... 
```

The set is considerably faster, especially for larger collections:

```
>>> compare(10)
ratio: 1.83441560587
>>> compare(100)
ratio: 4.4749036373
>>> compare(1000)
ratio: 21.4793493288
```

4 The Right Data Structure

```
>>> compare(10000)
ratio: 203.487480019
>>> compare(100000)
ratio: 1048.8407761
```

We did not measure the time it takes to convert the list into a set. So, let's define a modified function for the set that includes the creation of the set into the runtime measurement:

```
>>> def search_set_convert(n):
...     my_list = range(n)
...     start = timeit.default_timer()
...     my_set = set(my_list)
...     n in my_set
...     return timeit.default_timer() - start
...
```

we need a corresponding compare function:

```
>>> def compare_convert(n):
...     print 'ratio:', search_list(n) / search_set_convert(n)
...
```

Now the set is not faster anymore:

```
>>> compare_convert(10)
ratio: 0.456790136742
>>> compare_convert(100)
ratio: 0.316335542345
>>> compare_convert(1000)
ratio: 0.624656834843
>>> compare_convert(10000)
ratio: 0.405443366236
>>> compare_convert(100000)
ratio: 0.308628738218
>>> compare_convert(1000000)
ratio: 0.295318162219
```

If we need to search more than once, the overhead for creating the set gets relatively smaller. We write function that searches in our list several times:

```
>>> def search_list_multiple(n, m):
...     my_list = range(n)
...     start = timeit.default_timer()
...     for x in xrange(m):
...         n in my_list
...     return timeit.default_timer() - start
```

and do the same for our set:

```
>>> def search_set_multiple_convert(n, m):
...     my_list = range(n)
...     start = timeit.default_timer()
```

4 The Right Data Structure

```
...     my_set = set(my_list)
...     for x in xrange(m):
...         n in my_set
...     return timeit.default_timer() - start
```

We also need a new compare function:

```
>>> def compare_convert_multiple(n, m):
...     print 'ratio:', (search_list_multiple(n, m) /
...                     search_set_multiple_convert(n, m))
```

The set gets relatively faster with increasing collection size and number of searches.

```
>>> compare_convert_multiple(10, 1)
ratio: 0.774266745907
>>> compare_convert_multiple(10, 10)
ratio: 1.17802196759
>>> compare_convert_multiple(100, 10)
ratio: 2.99640026716
>>> compare_convert_multiple(100, 100)
ratio: 12.1363117596
>>> compare_convert_multiple(1000, 1000)
ratio: 39.478349851
>>> compare_convert_multiple(10, 1000)
ratio: 180.783828766
>>> compare_convert_multiple(10, 1000)
ratio: 3.81331204005
```

Let's assume we have two lists:

```
>>> list_a = list('abcdefg')
>>> list_a
['a', 'b', 'c', 'd', 'e', 'f', 'g']
>>> list_b = list('fghijklmnopq')
>>> list_b
['f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q']
```

and we would like to find out which letters are in both lists. A simple implementation would look like this:

```
>>> in_both = []
>>> for a in list_a:
...     if a in list_b:
...         in_both.append(a)
```

```
>>> in_both
['f', 'g']
```

This can be achieved in fewer lines and in most cases faster with sets:

```
>>> set_a = set(list_a)
>>> set_b = set(list_b)
```

4 The Right Data Structure

```
>>> set_a.intersection(set_b)
set(['g', 'f'])
```

Following the same method, we write a short performance test. First we write the function that uses lists:

```
>>> def intersect_list(n):
...     list_a = range(n)
...     list_b = range(n-3, 2 * n)
...     start = timeit.default_timer()
...     in_both = []
...     for a in list_a:
...         if a in list_b:
...             in_both.append(a)
...     run_time = timeit.default_timer() - start
...     return run_time, in_both
... 
```

and check if the results is what we expected:

```
>>> intersect_list(10)
(1.0189864042331465e-005, [7, 8, 9])
```

Now, we write a function for sets:

```
>>> def intersect_set(n):
...     set_a = set(range(n))
...     set_b = set(range(n-3, 2 * n))
...     start = timeit.default_timer()
...     in_both = set_a.intersection(set_b)
...     run_time = timeit.default_timer() - start
...     return run_time, in_both
... 
```

We are faster but the result of the intersection is the same:

```
>>> intersect_set(10)
(4.0926115616457537e-006, set([8, 9, 7]))
```

Finally, we write a comparison function in which we assert that both results are the same, and calculate the run time ratios.

```
>>> def compare_intersect(n):
...     list_time, list_result = intersect_list(n)
...     set_time, set_result = intersect_set(n)
...     assert set_result == set(list_result)
...     print 'ratio:', list_time / set_time
... 
```

Now we can compare both versions with lists and sets:

4.3 list vs. deque

```
>>> compare_intersect(10)
ratio: 2.75475854866
>>> compare_intersect(100)
ratio: 49.3294012578
>>> compare_intersect(1000)
ratio: 581.103479374
>>> compare_intersect(10000)
ratio: 7447.07128383
```

Note that the problem with the time for constructing the sets is not included here.

4.3 list vs. deque

For certain tasks we can use a deque instead of a list. A deque is a doubly linked list. This data structure allows faster insertion into the middle part. On the other hand, access of elements by index is slow.

So far we have instrumented our functions we want to test manually with timing code. It is far more elegant to move this timing code into its own, reusable module. In analogy to the decorator we wrote for profiling memory usage, we write one for speed in seconds and kilo stones:

```
# file: profile_speed.py

"""Profile the run time of a function with a decorator.
"""
import functools

import timeit #1

import pystone_converter #2

speed = {} #3

def profile_speed(function): #4
    """The decorator.
    """
    @functools.wraps(function)
    def _profile_speed(*args, **kwargs): #5
        """This replaces the original function.
        """
        start = timeit.default_timer() #6
        try:
            return function(*args, **kwargs) #7
        finally:
            # Will be executed *before* the return.
            run_time = timeit.default_timer() - start #8
            #9
            kstones = pystone_converter.kpystone_from_seconds(run_time)
            speed[function.__name__] = {'time': run_time, #10
                                       'kstones': kstones} #11

    return _profile_speed
```

We need the time module (#1) to measure the elapsed time. We also import our converter from seconds to pystones (#2). Again, we use a global dictionary to store our speed profiling results (#3). The decorator function takes function to be speed tested as argument (#4). The nested function takes positional and

4.3 list vs. deque

keyword arguments (#5) that will be supplied to the measured function. We record a time stamp for the start (#6) and call our function with arguments (#7). After this, we calculate the run time (#8) and convert it into kilo pystones (#9). Finally, we store the measured values in the global dictionary (#10) and return our nested function (#11).

Now we can use our module at the interactive prompt:

```
>>> import profile_speed
```

We decorate a function that takes a list and deletes several elements somewhere in the list by assigning an empty list to the range to be deleted:

```
>>> @profile_speed.profile_speed
... def remove_from_list(my_list, start, end):
...     my_list[start:end] = []
... 
```

Now we use a deque to do the same:

```
>>> @profile_speed.profile_speed
... def remove_from_deque(my_deque, start, end):
...     my_deque.rotate(-end)
...     for counter in range(end - start):
...         my_deque.pop()
...     my_deque.rotate(start)
... 
```

We rotate by `-end` to move the elements that need to be deleted to the end, call `pop` as many times as needed and rotate back by `start`.

Let's look at this rotating with a small example: We would like to achieve this:

```
>>> l = range(10)
>>> l[2:4] = []
>>> l
[0, 1, 4, 5, 6, 7, 8, 9]
```

We first make a deque:

```
>>> d = deque(range(10))
>>> d
deque([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

Now we rotate by the negative end index:

```
>>> d.rotate(-4)
>>> d
deque([4, 5, 6, 7, 8, 9, 0, 1, 2, 3])
```

We remove the last two elements:

4.3 list vs. deque

```
>>> d.pop()
3
>>> d.pop()
2
```

and rotate back in the desired order:

```
>>> d.rotate(2)
>>> d
deque([0, 1, 4, 5, 6, 7, 8, 9])
```

Now, let's test the speed of our implementations. We make a large list:

```
>>> my_list = range(int(1e6))
```

and import deque from the collections module:

```
>>> from collections import deque
```

We make a deque from our list:

```
>>> my_deque = deque(my_list)
```

Now we call both of our decorated functions:

```
>>> remove_from_list(my_list, 100, 105)
>>> remove_from_deque(my_deque, 100, 105)
```

The speed measuring results are in the global dictionary speed in profile_speed:

```
>>> profile_speed.speed['remove_from_list']
{'kstones': 0.05940467108987868, 'time': 0.0015446220713783987}
>>> profile_speed.speed['remove_from_deque']
{'kstones': 0.00090945420190496104, 'time': 2.3647349735256284e-005}
```

To be able to compare the results better, we calculate the ratio of both speeds:

```
>>> (profile_speed.speed['remove_from_list']['kstones'] /
...  profile_speed.speed['remove_from_deque']['kstones'])
71.706250305342934
```

Our deque is considerably faster than our list. But now we increase the range that is to be deleted:

```
>>> remove_from_list(my_list, 100, 1000)
>>> remove_from_deque(my_deque, 100, 1000)
```

And get a much smaller gain by using a deque:

```
>>> (profile_speed.speed['remove_from_list']['kstones'] /
...  profile_speed.speed['remove_from_deque']['kstones'])
```

4.4 dict vs. defaultdict

```
4.925948467147018
```

We make the range even larger:

```
>>> remove_from_list(my_list, 100, 10000)
>>> remove_from_deque(my_deque, 100, 10000)
```

Our list eventually becomes faster than the deque:

```
>>> (profile_speed.speed['remove_from_list']['kstones'] /
...  profile_speed.speed['remove_from_deque']['kstones'])
0.5219062068409327
```

4.4 dict vs. defaultdict

Since Python 2.5 there is new `defaultdict` in the module `collections`. This works similarly to the `defaultdict` method of dictionaries.

Let's assume we want to count how many of each letter are in the following sentence:

```
>>> s = 'Some letters appear several times in this text.'
```

We can do this in the standard way:

```
>>> d = {}
>>> for key in s:
...     d.setdefault(key, 0)
...     d[key] += 1
...
>>> d
{'a': 3, ' ': 7, 'e': 8, 'i': 3, 's': 4, 'm': 2,
 'l': 2, 'o': 1, 'n': 1, 'p': 2, 'S': 1, 'r': 3,
 't': 6, 'v': 1, 'x': 1, 'h': 1, '.': 1}
```

Or we can use the new `defaultdict`:

```
>>> dd = collections.defaultdict(int)
>>> for key in s:
...     dd[key] += 1
...
>>> dd
defaultdict(<type 'int'>, {'a': 3, ' ': 7, 'e': 8, 'i': 3, 's': 4, 'm': 2,
 'l': 2, 'o': 1, 'n': 1, 'p': 2, 'S': 1, 'r': 3, 't': 6, 'v': 1, 'x': 1,
 'h': 1, '.': 1})
>>>
```

Let's profile the speed differences. First, a function with our standard dictionary:

```
>>> @profile_speed.profile_speed
... def standard_dict(text):
...     d = {}
```

4.4 dict vs. defaultdict

```
...     for key in text:
...         d.setdefault(key, 0)
...         d[key] += 1
...
```

And now one for the defaultdict:

```
>>> import collections
>>> @profile_speed.profile_speed
... def default_dict(text):
...     dd = collections.defaultdict(int)
...     for key in text:
...         dd[key] += 1
...
```

We call them both with the same data:

```
>>> standard_dict(s)
>>> default_dict(s)
```

and compare the results:

```
>>> (profile_speed.speed['standard_dict']['kstones']
     profile_speed.speed['default_dict']['kstones'])
1.0524903876080238
```

There is not much difference between them: Therefore, we increase the size of our data:

```
>>> s = 'a' * int(1e6)
>>> standard_dict(s)
>>> default_dict(s)
```

and get a more than twofold speedup:

```
>>> (profile_speed.speed['standard_dict']['kstones'] /
     profile_speed.speed['default_dict']['kstones'])
2.3854284818433915
```

Let's look at different example from the Python documentation. We have this data structure:

```
>>> data = [('yellow', 1), ('blue', 2), ('yellow', 3), ('blue', 4), ('red', 1)]
```

Our goal is to produce a dictionary that groups all second tuple entries into a list:

```
>>> d.items()
[('blue', [2, 4]), ('red', [1]), ('yellow', [1, 3])]
```

Again, we define a decorated function for the two dictionary versions:

4.4 dict vs. defaultdict

```
>>> @profile_speed.profile_speed
... def default_dict_group(data):
...     dd = collections.defaultdict(list)
...     for key, value in data:
...         dd[key].append(value)
... 
```

```
>>> @profile_speed.profile_speed
... def standard_dict_group(data):
...     d = {}
...     for key, value in data:
...         d.setdefault(key, []).append(value)
... 
```

Call them:

```
>>> default_dict_group(data)
>>> standard_dict_group(data)
```

and look at the results:

```
>>> (profile_speed.speed['standard_dict_group']['kstones'] /
    profile_speed.speed['default_dict_group']['kstones'])
0.69018090107868191
```

The defaultdict seems to be slower. So let's increase the data size:

```
>>> data = data * 10000
>>> standard_dict_group(data)
>>> default_dict_group(data)
```

Now we are nearly twice as fast:

```
>>> (profile_speed.speed['standard_dict_group']['kstones'] /
    profile_speed.speed['default_dict_group']['kstones'])
1.9115965603608458
```

Making the data even larger makes things only slightly faster:

```
>>> data = data * 10
>>> standard_dict_group(data)
>>> default_dict_group(data)
```

```
>>> (profile_speed.speed['standard_dict_group']['kstones'] /
    profile_speed.speed['default_dict_group']['kstones'])
1.9823501285360818
```

Another increase by a factor of ten actually produces a less favorable ratio for the defaultdict:

```
>>> data = data * 10
>>> standard_dict_group(data)
>>> default_dict_group(data)
```

```
>>> (profile_speed.speed['standard_dict_group']['kstones'] /
     profile_speed.speed['default_dict_group']['kstones'])
1.8241023044794571
```

4.5 Big-O notation and Data Structures

Normally, you would want to reduce complexity of your program to make it faster. One frequently used measure for complexity is the so called **big-O** notation. The following table gives an overview of some notations along with a short description and some examples from Python.

Notation	Description	Python Examples
$O(1)$	constant time does not increase with size of data	<code>len(my_list)</code> , <code>len(my_dict)</code> , <code>my_list[i]</code> , <code>del my_dict[i]</code> , <code>x in dict</code> , <code>x in set</code> , <code>my_list.append(i)</code>
$O(n)$	linear time increase linearly with size of data	Loops on list, strings, dicts, sets, string methods, <code>x in my_list</code>
$O(n \log n)$	quasi linear time increases a little faster than linearly	<code>my_list.sort()</code>
$O(n^2)$	quadratic time increases four times for each doubling of data	nested loops
$O(n^3)$	cubic time increases four times for each doubling of data	nested nested loops
$O(n!)$	factorial	traveling sales man problem (not Python specific)

In general, using big-O notation we look only at the order of magnitude. Constant factors are neglected. So $O(3*n)$ and $O(20*n)$ are called $O(n)$. Therefore, $O(20*n)$ might be slower than $O(n^2)$ for very small n . But for large n the constant factor has very little influence.

Actually we have already compared several of these notations in our examples above. Let's look at some more comparisons of notations.

4.6 $O(1)$ vs. $O(n)$ vs. $O(n^2)$

We use our decorator from the module `profile_speed`:

```
>>> import profile_speed
```

We write a function that takes an iterable and reverses it into a list. Our first implementation uses the method `insert` to insert every item at the first position:

```
>>> @profile_speed.profile_speed
... def use_on(iterable):
...     result = []
...     for item in iterable:
...         result.insert(0, item)
```

4.5 Big-O notation and Data Structures

```
...     return result
...
```

Our second implementation uses `append` and reverse the list after all items are appended:

```
>>> @profile_speed.profile_speed
... def use_o1(iterable):
...     result = []
...     for item in iterable:
...         result.append(item)
...     result.reverse()
...     return result
...
```

Now we compare both functions in terms for runtime:

```
>>> def compare_on_o1(n):
...     r1 = use_on(range(n))
...     r2 = use_o1(range(n))
...     assert r1 == r2
...     print (profile_speed.speed['use_on']['kstones'] /
...           profile_speed.speed['use_o1']['kstones'])
...
>>> compare_on_o1(10)
1.6353049525
>>> compare_on_o1(100)
2.01816718953
>>> compare_on_o1(1000)
4.04768995537
>>> compare_on_o1(10000)
27.2673621812
>>> compare_on_o1(100000)
156.635364154
>>> compare_on_o1(int(1e6)) # this might take a while
2355.86619878
```

The speed differences are growing rapidly with increasing data sizes. The method `append` is $O(1)$ and `reverse` is $O(n)$. Even though `insert` is also $O(n)$ it is called n times whereas `reverse` is called only once. Because we loop over all items of our iterable, the first function is $O(n + n)$ but the second is $O(n^2)$. Putting this in numbers, we get:

```
>>> for x in [10, 100, 1000, 10000, 100000, 1000000]:
...     print x * x / (x + x)
...
5
50
500
5000
50000
500000
```

Of course instead of appending to a new list we can just convert the iterable into a list and reverse it:

4.7 Exercises

```
>>> @profile_speed.profile_speed
... def use_list(iterable):
...     result = list(iterable)
...     result.reverse()
...     return result
...
```

Now we can compare both implementations that have the same big-O notation:

```
>>> def compare_ol_list(n):
...     r1 = use_list(range(n))
...     r2 = use_ol(range(n))
...     assert r1 == r2
...     speed = profile_speed.speed
...     print (speed['use_ol']['kstones'] /
...            speed['use_list']['kstones'])
...
```

```
>>> compare_ol_list(10)
1.24255753768
>>> compare_ol_list(100)
4.39513352799
>>> compare_ol_list(1000)
23.1481811661
>>> compare_ol_list(10000)
54.2245839131
>>> compare_ol_list(100000)
53.132471733
>>> compare_ol_list(1000000)
29.8124806601
```

Even though the big-O notation is the same, the `list` version is up to 50 times faster.

4.7 Exercises

1. Write a test program that searches the last number in a long list. Use `item` in `long_list` and `item` in `set(long_list)`. Perform this search 10 and more times. Compare the run times. Hint: You can use `timeit.default_timer()` to get the time since the last call to this function. Alternatively, you can use the module `timeit` or the function `measureRunTime` which you can find in the `examples` directory in the subdirectory `modules`.

5 Caching

5.1 Reuse before You Recalculate

If you find yourself calling the same function with the same arguments many time then caching might help to improve the performance of your program. Instead of doing an expensive calculation, database query, or rendering again and over again, caching just reuses the results of former function calls. Depending on whether the results will be the same for every call to the same function with the same arguments or if the result might change over time, we talk about deterministic or non-deterministic caching. An example for deterministic caching would be numerical calculations that should always produce the same result for the same input. Caching of database queries is non-deterministic because the database content might change. So after some timeout period the query has to be done anew.

All of the following examples are based on [ZIAD2008].

5.2 Deterministic caching

The first thing we need to do, if we want to cache function results, is to uniquely identify the function we want to call:

```
# file: get_key.py
# based on Ziade 2008

"""Generate a unique key for a function and its arguments.
"""

def get_key(function, *args, **kw):                                     #1
    """Make key from module and function names as well as arguments.
    """
    key = '%s.%s:' % (function.__module__,                             #2
                      function.__name__)
    hash_args = [str(arg) for arg in args]                             #3
    hash_kw = ['%s:%s' % (k, str(v))                                   #4
               for k, v in kw.items()]
    return '%s::%s::%s' % (key, hash_args, hash_kw)                  #5
```

The function `get_key` takes a function and its positional and keyword arguments (#1). We extract the module name and function name from the function (#2). Now we convert all positional arguments into a list of strings (#3). We convert the keyword arguments into a list of strings using the keys and the string representation of the values (#4). Finally, we return a string that consists of the three strings we have assembled so far (#5).

Now we use our function for a decorator to memoize (a term for the kind of caching we perform) previously calculated results:

```
# file: cache_deterministic.py
# form Ziade 2008

"""Example for a deterministic cache
"""

import functools
```

5 Caching

```
from get_key import get_key #1

cache = {} #2

def memoize_deterministic(get_key=get_key, cache=cache): #3
    """Parameterized decorator for memoizing.
    """

    def _memoize(function): #4
        """This takes the function.
        """

        @functools.wraps(function)
        def __memoize(*args, **kw): #5
            """This replaces the original function.
            """

            key = get_key(function, *args, **kw) #6
            try:
                return cache[key] #7
            except KeyError: #8
                value = function(*args, **kw) #9
                cache[key] = value #10
                return value

            return __memoize
        return __memoize
```

We use our function `get_key` (#1) and define a global dictionary that will be used to store pre-calculated data (#2). Our decorator takes the function and the dictionary as arguments (#3). This allows us to use other functions to retrieve a key and other caches possibly data dictionary-like data stores such as `shelve`. The second level function takes the function that is to be called as argument (#4). The third level function takes the arguments (#5). Now we retrieve our key (#6) and try to access the result from our cache (#7). If the key is not in the cache, we call our function (#8), store the result in the cache (#9) and return the result (#10).

Let's try how it works. We import the time modul and our module with the decorator:

```
>>> import time
>>> import cache_deterministic
```

We define a new function that adds to numbers and is decorated:

```
>>> @cache_deterministic.memoize_deterministic()
... def add(a, b):
...     time.sleep(2)
...     return a + b
... 
```

We simulate some heavy calculations by delaying everything for two seconds with `sleep`. Let's call function:

5.3 Non-deterministic caching

```
>>> add(2, 2)
4
```

This took about two seconds. Do it again:

```
>>> add(2, 2)
4
```

Now the return is immediate.

Again:

```
>>> add(3, 3)
6
```

Two seconds delay. But now:

```
>>> add(3, 3)
```

Instantaneous response.

5.3 Non-deterministic caching

For non-deterministic caching, we use an age that the computed value should not exceed:

```
# file: cache_non_deterministic.py
# form Ziade 2008

"""Example for a cache that expires.
"""

import functools
import time

from get_key import get_key

cache = {}

def memoize_non_deterministic(get_key=get_key, storage=cache,
                              age=0): #1
    """Parameterized decorator that takes an expiration age.
    """

    def _memoize(function):
        """This takes the function.
        """

        @functools.wraps(function)
        def __memoize(*args, **kw):
            """This replaces the original function.
            """

            key = get_key(function, *args, **kw)
            try:
```

5.3 Non-deterministic caching

```
        value_age, value = storage[key]                #2
        deprecated = (age != 0 and                    #3
                      (value_age + age) < time.time())
    except KeyError:
        deprecated = True                             #4
    if not deprecated:
        return value                                  #5
    storage[key] = time.time(), function(*args, **kw)  #6
    return storage[key][1]                             #7
return __memoize
return __memoize
```

This decorator is a variation of the deterministic one above. We can supply an age (#1). The value will be recalculated if this age is exceeded. We retrieve an age and a value from our cache (#2). The value will be deprecated, i.e. recalculated if we provide a non-zero age and the old age plus the specified age are smaller than the current time (#3). Note: This means, if you provide no age or an age of zero, the cache will never expire. The value will also be calculated if the key was not found (#4). We return the value if it is still valid (#5). Otherwise, we recalculate it and store it together with current time in the cache (#6) and return the freshly calculated value (#7).

Let's see how this works. We import our non-deterministic cache:

```
>>> import cache_non_deterministic
```

and define a new function with a maximum cache age of 5 seconds:

```
>>> @cache_non_deterministic.memoize_non_deterministic(age=5)
... def add2(a, b):
...     time.sleep(2)
...     return a + b
... 
```

The first call takes about two seconds:

```
>>> add2(2, 2)
4
```

Immediately after this we do it again and get the response without delay:

```
>>> add2(2, 2)
4
```

Now we wait for at least 5 seconds ... and do it again:

```
>>> add2(2, 2)
4
```


This took again two seconds because the cache was not used and the value was recalculated due to the time since the last call being greater than five seconds.

5.4 Memcached

[Memcached](#) ⁶ is a caching server that is primarily used to speed up database based dynamic web pages. It is very fast, trades RAM for speed, and is very powerful. We don't have time to look at it here. There are several ways to use Memcached from Python. Also the probably most popular web framework Django uses Memcached ([Djangos cache](#) ⁷).

6 The Example

We will use the example of a very simple Monte Carlo simulation to determine pi for all our next topics. The approximate calculation of pi can be done by finding the ratio of points that lie within and outside a circle but inside a square with side length of the circles radius as shown in the figure. We look only at the upper right quarter of the circle. By doing so, we can place the origin of our coordinate system in the center of the circle and still have only positive coordinates.



The area of the quarter of the circle is calculated as:

$$A_c = \frac{\pi r^2}{4}$$

The area of the square is:

$$A_s = r^2$$

So their ratio is:

$$\frac{A_c}{A_s} = \frac{\frac{\pi r^2}{4}}{r^2}$$

Through conversion we get:

$$\pi = 4 \frac{A_c}{A_s}$$

Assuming an infinite number of points the ratio of points may be used instead of the areas, so:

$$\pi = 4 \frac{\text{number of points inside}}{\text{total number of points}}$$

7 Testing Speed

Testing the speed of a program or its portions is a difficult task because many external circumstances can influence the result. Python offers the module `timeit` that is designed to check small pieces of code at the commandline

```
$ python -mtimeit 'x=0' 'x+=1'
10000000 loops, best of 3: 0.132 usec per loop
```

or from within a module. We use `timeit` in the module `measure_time.py` to run all our test functions that calculate pi with the Monte Carlo method.

```
# file: measure_time.py

"""Measure the run time of functions.
"""

from timeit import Timer #1

def measure_run_time(total, names, number=1): #2
    """Measure the run times of all functions given with `names`.
    """

    def timeit_runner(name): #3
        """Time one function.
        """
        timer = Timer('%s(total)' % name,
                       'from __main__ import %s\ntotal=%d' #4
                       % (name, total)) #5
        return timer.timeit(number), name #5

    results = [] #6
    for name in names: #7
        results.append(timeit_runner(name)) #7
    results.sort() #8
    length = max(len(name) for name in (names)) #9
    format1 = '%s-%ds' % length
    header = (format1 + '%s%10s') % ('Function',
                                   'Time'.center(11), 'Ratio')

    print
    print header
    print '=' * len(header)
    for result in results:
        ratio = result[0] / results[0][0]
        print (format1 + '%9.4f s%10.2f') % (result[1],
                                             result[0], ratio)

if __name__ == '__main__':

    import time

    def func1(total):
```

7 Testing Speed

```
    """Test function that waits for `total` seconds.
    """
    time.sleep(total)

def func2(total):
    """Test function that waits for `total` * 2 seconds.
    """
    time.sleep(total * 2)

def test():
    """Check if it works.
    """
    measure_run_time(1, ['func1', 'func2'])

test()
```

We import the class `Timer` from the module `timeit` (#1). The function `measure_run_time` takes the number of iterations for the Monte Carlo calculation `n`, the names of the functions to run `names` and the number of runs. It averages run time speed over `number` defaulting to 1 (#2).

Within this function we have the helper function `timeit_runner` (#3) that takes only the name of the function. It makes an instance of `Timer` (#4) and runs the function the specified number of times. It returns the runtime and the name of the function (#5).

The rest of `measure_run_time` collects all results in a list (#6, #7), sorts them fastest first (#8), and prints them to the screen in a formatted table (#9 and the following lines).

The output looks like this:

measure_time.py

Function	Time	Ratio
func1	0.9993 s	1.00
func2	1.9999 s	2.00

8 Pure Python

Our first algorithm is plain python:

```
# file: plain_pi.py

"""Calculating pi with Monte Carlo.
"""

import math #1
import random
import sys

# make `import measure_time` work
sys.path.append('../measuring')
import measure_time #2

def pi_plain(total): #3
    """Calculate pi with `total` hits.
    """
    count_inside = 0 #4
    for _ in xrange(total): #5
        x = random.random() #6
        y = random.random() #7
        dist = math.sqrt(x * x + y * y) #8
        if dist < 1: #9
            count_inside += 1 #9
    return 4.0 * count_inside / total #10

if __name__ == '__main__': #11

    def test():
        """Check if it works.
        """
        print 'pi:', pi_plain(int(1e5)) #12
        names = ['pi_plain'] #13
        total = int(1e5) #14
        repeat = 5 #15
        measure_time.measure_run_time(total, names, repeat) #16

    test()
```

We first import the modules `random` and `math` (#1), and then our module `measure_time` explained above (#2). The function `pi_plain` takes the number of intended iterations as argument (#3). First, we set the counter for the number of points inside the circle to zero (#4). Then we loop over the range from 0 to `total` (#5). The `x` and `y` coordinates are random numbers in the range (0,1) (#6, #7). Following Pythagoras, the distance `dist` from the origin is the square root of the sum of the squares of the coordinates (#8). If this distance is smaller than 1, we count it as inside the circle (#9). We return the ratio of the inside and total count multiplied by four (#10) after the loop has terminated. Using the Python for the main program (#11), we print the result of the function call to the screen (#12). Now we start our performance test using a list with names of the functions to test. In this case it is only one (#13). We want 100,000 iterations. Because it is easy to loose count with all the zeros, we use a notation frequently used in science `1e5`. Since this is a float and an integer is needed, we convert it (#14). We want the time to be

8 Pure Python

measured five times (#15). Finally, we start the measurement (#16). The result should be (with different numbers for `pi` and `Time`):

```
$ plain_pi.py
pi: 3.142032

Function      Time          Ratio
=====
pi_plain      0.8327 s          1.00
```

Since we have only one function, the ratio is 1.0. This ratio is always the one for the fastest function and bigger for the slower ones. A ratio of 2.0 means it takes twice as long as the fastest.

Now we make some small changes to the program. The module `math` in the Python standard library has the function `hypot(x, y)`, which is equivalent to the `sqrt(x*x + y*y)`. Assuming it is faster, we modify the program:

```
# file: math_pi.py

"""Variation on `pi_plain` but using `hypot`
from the `math` module.
"""

import math
import random

def pi_math(total):                                     #1
    """Pi calculation using hypot.
    """
    count_inside = 0
    for _ in range(total):
        dist = math.hypot(random.random(),             #2
                           random.random())
        if dist < 1:
            count_inside += 1
    return 4.0 * count_inside / total

if __name__ == '__main__':

    import plain_pi
    import measure_time
    pi_plain = plain_pi.pi_plain

    def test():
        """Check if it works.
        """

        names = ['pi_plain', 'pi_math']
        total = int(1e5)
        repeat = 5
        measure_time.measure_run_time(total, names, repeat)

    test()
```

8.1 Exercise

This is very similar to the first program, only the name of the function (#1) and the method to calculate the distance (#2) have changed. Now, we would like to test the performance and compare it with that of the first function. Because we want all functions in separate files, we write a small testing module. The rest of the source is identical with our first program. The output of the program should be like this:

```
$ math_pi.py

Function      Time          Ratio
=====
pi_plain      0.8373 s        1.00
pi_math       1.1272 s        1.35
```

Obviously, the `hypot` is slightly slower than our more detailed calculation. This result is not necessarily expected.

8.1 Exercise

1. Use the function `pi_plain` and replace `range` with `xrange`. Time it, applying the same procedure shown above. Compare the results. Vary the numbers `n` and `m`. How do results change?
2. Use the power function `pow()` and `**` to calculate the squares, i.e. `pow(x, 2)` and `x**2` instead of `x * x`.
3. Use local names instead of global variables. For instances, avoid `module.function()` by making `function` local before the loop.
4. Suggest more changes to improve performance.

9 Meet Psyco, the JIT

Just-in-time compilation (JIT) is a method to bring the speed of compiled languages to dynamic languages such as Python. [Psyco](#)⁸ is a tool that can help to improve Python performance considerably. Depending on the type of code, it is possible to achieve several-fold speedups.

No New Development but Still Useful at Times

Psyco is **not** developed any further. The developers moved on to PyPy, our next topic. So Psyco is only available for the 32-bit PC architecture and Python ≤ 2.6 . On the other hand, there are quite a few old Python installations around that can profit from Psyco and the effort to apply Psyco is relatively small.

For small applications the whole program can be optimized by adding the following two lines at the beginning of your main module:

```
import psyco
psyco.full()
```

It might be desirable to optimize only parts of the application. To find the parts that profit the most from optimization, you should do profiling with Psyco:

```
import psyco
psyco.log()
psyco.profile()
```

Psyco will do profiling to find functions that are more time consuming than others. The results are written to a log file with the name of your module with the extension `log-psyco`. The functions found to consume the most CPU time are compiled. You can fine tune this process by supplying arguments. You can also select the functions to optimize with Psyco by hand:

```
import psyco
psyco.bind(pi_plain)
```

If you want to compare the run time of a function and its psyco-compiled version in one program, you can use Psycos `proxy` function. It returns a proxy to the original function which remains unchanged:

```
import psyco
fast_func = psyco.proxy(slow_func)
slow_func()
fast_func()
```

9.1 Exercises

1. Apply Psyco to all functions implemented so far. Time them, applying the same procedure shown above. Compare the results.
2. Try different methods such as `full`, `profile` or `bind` and compare the results for different numbers for `n` and `m`.

10 PyPy: New, Fast, Maturing

PyPy is an implementation of Python 2.7. that is faster than the standard CPython for many common tasks. In addition it supports sandboxing to run untrusted code and stackless support besides other features.

As long as the code is pure Python and runs with Python 2.7, it should run with PyPy. C code can be accessed via ctypes and there is active development to use Python C extensions.

Let's just run our code in PyPy and compare the run time to that of CPython for 1e6 throws:

```
python plain_pi.py
pi: 3.1426

Function      Time          Ratio
=====
pi_plain      5.0093 s        1.00

pypy plain_pi.py
pi: 3.144164

Function      Time          Ratio
=====
pi_plain      0.6998 s        1.00
```

This gives us a 7 times faster execution without doing anything except using *pypy* instead of *python* at the command line.

The `math_pi.py` example is about 6 times faster:

```
python math_pi.py

Function      Time          Ratio
=====
pi_math       4.0295 s        1.00

pypy math_pi.py

Function      Time          Ratio
=====
pi_math       0.6784 s        1.00
```

10.1 Exercises

1. Use PyPy for the pure Python algorithms for calculating pi. Vary the number of iterations. Compare the results with runs with CPython.

11 NumPy for Numeric Arrays

Python is used by many scientists who often have to deal with large amounts of data in form of multidimensional arrays. NumPy⁹ is a library implemented in C that brings matlab-like numerical processing to Python. We apply NumPy to our example:

```
# file: numpy_pi.py
"""Calculating pi with Monte Carlo Method and NumPy.
"""

import numpy #1

def pi_numpy(total): #2
    """Compute pi.
    """
    x = numpy.random.rand(total) #3
    y = numpy.random.rand(total) #4
    dist = numpy.sqrt(x * x + y * y) #5
    count_inside = len(numpy.where(dist < 1)[0]) #6
    return 4.0 * count_inside / total

if __name__ == '__main__':
    def test():
        """Time the execution.
        """
        import timeit
        start = timeit.default_timer()
        print 'pi:', pi_numpy(int(1e6))
        print 'run time', timeit.default_timer() - start
    test()
```

After we import `numpy` (#1), we define the function `numpy_pi` that again takes the number of iterations as argument `total` (#2). The function `numpy.random.rand` generates an array of random numbers in the range (0,1). So `x` and `y` hold such arrays (#3, #4). `dist` is the array of distances computed from `x` and `y` (#5). `where` returns an array with the values for which the condition `dist < 1` is true. Its length is the number of hits inside the circle (#6). The return value is a tuple with one element. Therefore, we need to index with zero `[0]`.

11.1 Exercise

Time `numpy_pi` applying the same procedure shown above. Compare the results. How do they depend on `total` and repeat?

12 NumPyPy

NumPy is widely used by scientists and engineers. Since this group of people would greatly profit from using PyPy, there is a recent effort to rewrite NumPy in PyPy: NumPyPy. This is not an interfacing effort to connect to the existing C code base of NumPy but rather a rewrite of NumPy in PyPy.

The task is not complete yet and there are some missing features. Let's try to use NumPyPy for our example. It turns out that in the current release of PyPy (1.8) the module for the random numbers and the `where` function are still missing. In addition there is still lots of work being done in speeding array construction. So NumPyPy is not ready yet to be compared to NumPy.

13 Using Multiple CPUs with multiprocessing

The future is parallel. CPUs don't perform so much faster anymore. The trend is towards multiple processors. Hyper threading, dual core, quad core and other terms are ubiquitous. Python has the GIL - the global interpreter lock - that prevents threads in a Python program to take advantage of more than one CPU. Since Python 2.6, the module `multiprocessing` is a part of the standard library. For older versions of Python there is the module `pyprocessing` that needs to be installed separately but does the same job.

Let's try to speed up our calculation with going parallel:

```
# file: processes_pi.py

"""Calculation of pi with Monte Carlo and multiprocessing.
"""

import math
try:
    from multiprocessing import Process, Queue          #1
except ImportError:
    from processing import Process, Queue              #2
import random
import timeit

def count_inside(total):                                #3
    """Count the hits inside the circle.

    This function will be run multiple times in different
    processes.
    """
    inside_count = 0
    for _ in xrange(total):
        x = random.random()
        y = random.random()
        dist = math.sqrt(x * x + y * y)
        if dist < 1:
            inside_count += 1
    return inside_count

def calc_pi(total):                                    #4
    """Determine pi _without_ multiprocessing as refrence.
    """
    return 4 * count_inside(total) / float(total)

def count_inside_process(queue, total):                 #5
    """Count the hits inside the circle in a separte process.
    """
    queue.put(count_inside(total))
```


13 Using Multiple CPUs with multiprocessing

```
def calc_pi_processes(total, process_count): #6
    """Calculate pi spread of several processes.

    We need to split the task into sub tasks before we can hand
    them to other processes.
    """
    min_n = total // process_count #7
    counters = [min_n] * process_count #8
    reminder = total % process_count #9
    for count in xrange(reminder): #10
        counters[count] += 1 #11
    queues_processes = [] #12
    for counter in counters: #13
        queue = Queue() #14
        process = Process(target=count_inside_process,
                          args=(queue, counter)) #15
        process.start() #16
        queues_processes.append((queue, process)) #17
    inside_count = sum(process[0].get() for process
                       in queues_processes) #18
    for queue_process in queues_processes: #19
        queue_process[1].join() #20
    return 4 * inside_count / float(total) #21
```

We need `Process` and `Queue` from `multiprocessing` (#1) or from `pyprocessing` (#2) if we work with pre 2.6 versions of Python. We use our method for counting the points inside the circle (#3). We use the non-parallel calculation of pi as reference (#4). We write a wrapper around this function that outs the calculated result into a queue (#5). Our function for calculating pi in parallel takes the number of points to investigate and the number of processes the calculation should be performed on as arguments (#6). We determine which share each process should calculate. The goal is to distribute the load evenly. Therefore, divided all tries (n) by the number of processes (#7) and put this result in list (#8). We calculate the reminder (#9). If there is a reminder, we have to add the remaining numbers starting from the first list element (#10 and #11).

We use a list for storing our queues and processes (#12). Going through all counters (#13), we create a new queue (#14) and start a new process (#15) with our function `count_inside_process` that will be called in each process with queue and the counter as arguments. Every step gets started (#16) and we append the process and the queues to our list (#17) to make them accessible for later use. We get the calculated result from each queue and sum them up (#18). Now we go through all processes (#19) and wait for them to finish (#20). Finally, we can calculate our pi using the result (#21).

```
if __name__ == '__main__':
    def test():
        """Check if it works.
        """
        process_count = 2 #22
        total = int(1e3) #23
        print 'number of tries: %2.0e' % total #24

        start = timeit.default_timer() #25
        print 'pi:', calc_pi(total)
        one_time = timeit.default_timer() - start
        print 'run time one process:', one_time
```

13 Using Multiple CPUs with multiprocessing

```
start = timeit.default_timer()                                #26
pi_value = calc_pi_processes(total, process_count)            #27
print 'pi', pi_value
two_time = timeit.default_timer() - start
print 'run time %d processes:' % process_count, two_time

print 'ratio:', no_time / two_time                            #28
print 'diff:', two_time - no_time

test()
```

We test our implementation with two processes (#22) and a number of tries (#23). Now we print the number of tries on the screen (#24). The test code for our non-process version is straight forward (#25).

We do the same thing for our process version (#26). The only difference is to use the wrapper function (#27). Finally, we show time ratios and time difference between the process and non-process version (#28).

On a computer with a dual core we get the following results:

```
processes_pi.py
number of tries: 1e+003
pi: 3.196
run time no processes: 0.00268888923
pi 3.14
run time 2 processes: 0.929015838605
ratio: 0.00289434164478
diff: 0.926326949375
```

```
processes_pi.py
number of tries: 1e+006
pi: 3.139364
run time no processes: 1.18375324238
pi 3.146548
run time 2 processes: 1.61413168433
ratio: 0.73336844439
diff: 0.430378441953
```

```
processes_pi.py
number of tries: 1e+007
pi: 3.1408952
run time no processes: 11.9309476484
pi 3.1422716
run time 2 processes: 7.79187159262
ratio: 1.53120434629
diff: -4.13907605576
```

```
processes_pi.py
number of tries: 1e+008
pi: 3.14183468
run time no processes: 118.451061238
pi 3.14122748
run time 2 processes: 71.0664637545
```

13 Using Multiple CPUs with multiprocessing

```
ratio: 1.66676453253
diff: -47.3845974838
```

We get a speedup of a factor of up to 1.6 for large numbers of iterations. On a different machine with an i5 processor that gives 4 cores according to `cpu_count`, the best speedup comes with 10 million iterations and three processes with more than double the speed of a single process:

```
number of tries: 1e+07
pi: 3.1412816
run time no processes: 8.22669792175
pi 3.1421968
run time 4 processes: 3.52357602119
ratio: 2.33475817529
diff: -4.70312190056
```

There are more possibilities to work with multiprocessing. For our use case, a pool is a good alternative:

```
# file pool_pi.py

"""Multiprocessing with a pool of workers.
"""

import math
try:
    from multiprocessing import Pool, cpu_count          #1
except ImportError:
    from processing import Pool, cpu_count               #2
import random
import timeit

def count_inside(total):
    """Count the hits inside the circle.

    This function will be run multiple times in different
    processes.
    """
    inside_count = 0
    for _ in xrange(total):
        x = random.random()
        y = random.random()
        dist = math.sqrt(x * x + y * y)
        if dist < 1:
            inside_count += 1
    return inside_count

def calc_pi(total):
    """Determine pi _without_ multiprocessing as reference.
    """
    return 4 * count_inside(total) / float(total)
```

13 Using Multiple CPUs with multiprocessing

```
def calc_pi_workers(total, workers=None): #3
    """Calculate pi spread of several processes with workers.

    We need to split the task into sub tasks before we can hand
    them to other process that run workers
    """
    if not workers:
        workers = cpu_count() #4
    min_n = total // workers
    counters = [min_n] * workers
    reminder = total % workers
    for count in xrange(reminder):
        counters[count] += 1
    pool = Pool(processes=workers) #5
    results = [pool.apply_async(count_inside, (counter,))
               for counter in counters] #6
    inside_count = sum(result.get() for result in results) #7
    return 4 * inside_count / float(total) #8

if __name__ == '__main__':

    def test():
        """Check if it works.
        """
        workers = 2

        total = int(1e3)

        print 'number of tries: %2.0e' % total
        start = timeit.default_timer()
        print 'pi:', calc_pi(total)
        no_time = timeit.default_timer() - start
        print 'run time no workers:', no_time

        start = timeit.default_timer()
        pi_value = calc_pi_workers(total, workers)
        print 'pi', pi_value
        two_time = timeit.default_timer() - start
        print 'run time %d workers:' % workers, two_time
        print 'ratio:', no_time / two_time
        print 'diff:', two_time - no_time

    test()
```

There are only a few changes compared to the version above. We need to import `Pool` and `cpuCount` from `multiprocessing` (#1) or `processing` (#2). Our function that will run in parallel now takes the number of workers as second argument (#3). If we don't provide the number of workers, we use as many workers as available CPUs are determined by `multiprocessing` (#4). We create a new pool (#5) with specified number of workers. We call `apply_async` for each worker and store all results (#6). Now we get the values of the results and sum them up (#7). Finally, we calculate our pi as usual (#8).

Again, on a dual core machine the parallel implementation is up to 1.8 times faster than the one-process version:

13 Using Multiple CPUs with multiprocessing

```
pool_pi.py
number of tries: 1e+003
pi: 3.208
run time no workers: 0.00241594951314
pi 3.084
run time 2 workers: 0.734127661477
ratio: 0.0032909119761
diff: 0.731711711963
```

```
pool_pi.py
number of tries: 1e+006
pi: 3.141976
run time no workers: 1.19714572662
pi 3.14252
run time 2 workers: 1.55787928354
ratio: 0.768445757813
diff: 0.360733556919
```

```
pool_pi.py
number of tries: 1e+007
pi: 3.1412288
run time no workers: 11.8663374814
pi 3.1407312
run time 2 workers: 7.29266065939
ratio: 1.62716161298
diff: -4.57367682205
```

```
pool_pi.py
number of tries: 1e+008
pi: 3.1413308
run time no workers: 119.62180333
pi 3.141572
run time 2 workers: 66.502250705
ratio: 1.79876323074
diff: -53.1195526247
```

14 Combination of Optimization Strategies

We can combine several of the techniques we used. Let's use NumPy in conjunction with the parallel approach:

```
#file: pool_numpy_pi.py

import math
try:
    from multiprocessing import Pool, cpuCount
except ImportError:
    from processing import Pool, cpuCount
import random
import time

import numpy

def count_inside(n):
    x = numpy.random.rand(n)
    y = numpy.random.rand(n)
    d = numpy.sqrt(x*x + y*y)
    return numpy.sum(numpy.where(d < 1, 1, 0))

def calc_pi(n):
    return 4 * count_inside(n) / float(n)

def calc_pi_workers_numpy(n, workers=None):
    if not workers:
        workers = cpuCount()
    min_n = n // workers
    counters = [min_n] * workers
    reminder = n % workers
    for m in xrange(reminder):
        counters[m] += 1
    pool = Pool(processes=workers)
    results = [pool.apply_async(count_inside, [counter])
               for counter in counters]
    inside_count = sum(result.get() for result in results)
    return 4 * inside_count / float(n)

if __name__ == '__main__':
    import time

    workers = 2

    n = int(4e7)

    print 'number of tries: %2.0e' % n
    start = time.clock()
    print 'pi:', calc_pi(n)
    no_time = time.clock() - start
    print 'run time no workers:', no_time

    start = time.clock()
```

```

pi = calc_pi_workers_numpy(n, workers)
print 'pi', pi
two_time = time.clock() - start
print 'run time %d workers:' % workers, two_time
print 'ratio:', no_time/two_time
print 'diff:', two_time - no_time

```

The only difference to `pool_pi.py` is the use of NumPy in `count_inside` (#1). We just return the number of points inside the circle (#2).

Now we measure the performance of all our implementations:

```

# file: montecarlo_pi.py

#!/usr/bin/env python

"""Test runner for all implementations.
"""

import sys

sys.path.append('../pi') #1
sys.path.append('../multiprocessing')

has_psyco = False
try:
    import psyco #2
    has_psyco = True
except ImportError:
    print 'No psyco found doing tests without it.'

has_numpy = False
try:
    import numpy #3
    has_numpy = True
except ImportError:
    print 'No numpy found doing tests without it.'

import measure_time #4

from plain_pi import pi_plain #5
from math_pi import pi_math
from pool_pi import calc_pi_workers
if has_numpy:
    from numpy_pi import pi_numpy
    from pool_numpy_pi import calc_pi_workers_numpy

if has_psyco:
    psyco_pi_plain = psyco.proxy(pi_plain) #6
    psyco_pi_math = psyco.proxy(pi_math)
    psyco_calc_pi_workers = psyco.proxy(calc_pi_workers)
    if has_numpy:

```

14 Combination of Optimization Strategies

```
psyco_pi_numpy = psyco.proxy(pi_numpy)
psyco_calc_pi_workers_numpy = psyco.proxy(calc_pi_workers_numpy)

def main():
    """Run all tests that could be found.
    """
    total = int(float(sys.argv[1])) #7
    repeat = int(sys.argv[2]) #8
    names = ['pi_plain', 'pi_math', 'calc_pi_workers'] #9
    if has_numpy:
        names.extend(['pi_numpy', 'calc_pi_workers_numpy']) #10
    if has_psyco:
        names.extend(['psyco_pi_plain', 'psyco_pi_math', #11
                     'psyco_calc_pi_workers'])
        if has_numpy:
            names.extend(['psyco_pi_numpy', #12
                          'psyco_calc_pi_workers_numpy']) #12
    measure_time.measure_run_time(total, names, repeat) #13

if __name__ == '__main__':
    main()
```

We add the paths that hold our libraries to `sys.path` so that we can import them (#1). We try to import `psyco` (#2) and `numpy` (#3), but their use is optional because `Psyco` is only available for 32-bit PC system with Python <= 2.6 and `NumPy` is still incomplete for `PyPy`. We now import our run time measuring module (#4) as well as all our modules we want to test (#5). We bind new function with `Psyco` only if it is available (#6).

We use two command line arguments to stir our test. The number of iterations `total` (#7) and the number how often the test are supposed to run `repeat` (#8). We use a list of function names (#9) that we eventually hand over to our measurement function along with `total` and `repeat` (#13). Depending on the availability of `NumPy` (10, 12) and `Psyco` (#12), we add more function names.

Finally, we compare the run times of our work:

```
montecarlo_pi.py 1e4 1
```

Function	Time	Ratio
=====	=====	=====
pi_numpy	0.0023 s	1.00
psyco_pi_numpy	0.0023 s	1.00
psyco_pi_plain	0.0032 s	1.40
psyco_pi_math	0.0064 s	2.80
pi_plain	0.0121 s	5.27
pi_math	0.0146 s	6.40
calc_pi_workers_numpy	1.1274 s	493.27
psyco_calc_pi_workers_numpy	1.1309 s	494.82
psyco_calc_pi_workers	1.1472 s	501.94
calc_pi_workers	1.2358 s	540.73

The parallel computations are very much slower than all others if we use only 1e4 iterations.

14.1 Exercise

Using three orders of magnitude more iterations gives a different result:

```
montecarlo_pi.py 1e7 1
```

Function	Time	Ratio
=====		
psyco_pi_numpy	1.8983 s	1.00
pi_numpy	2.1135 s	1.11
psyco_calc_pi_workers_numpy	2.4557 s	1.29
calc_pi_workers_numpy	2.4921 s	1.31
psyco_pi_plain	2.8003 s	1.48
psyco_pi_math	6.1471 s	3.24
psyco_calc_pi_workers	7.7396 s	4.08
calc_pi_workers	7.8426 s	4.13
pi_plain	12.0987 s	6.37
pi_math	14.7275 s	7.76

Doubling the iterations makes the parallel versions the fastest:

```
montecarlo_pi.py 2e7 2
```

Function	Time	Ratio
=====		
calc_pi_workers_numpy	3.4627 s	1.00
psyco_calc_pi_workers_numpy	3.4844 s	1.01
psyco_pi_numpy	3.8041 s	1.10
pi_numpy	4.2014 s	1.21
psyco_pi_plain	5.6769 s	1.64
psyco_pi_math	12.4210 s	3.59
psyco_calc_pi_workers	14.1498 s	4.09
calc_pi_workers	14.1824 s	4.10
pi_plain	24.8832 s	7.19
pi_math	29.5806 s	8.54

Your results may look somewhat different, depending on your compiler and your platform.

In conclusion, the fastest are the parallel NumPy versions followed by the serial plain NumPy versions. Applying psyco to NumPy has little effect. `psyco_pi_plain` is next in the ranking. Here psyco does good job. The parallel version (no NumPy) seems not to profit from psyco. The slowest implementation is about 8 times lower than the fastest.

These numbers are only valid for this very example. Other programs that may involve more complicated algorithms and data structures are likely to exhibit a different ranking of the versions. In particular, if we can make use of Python built-in data structures, such as dictionary, the performance of pure Python can be much better. This is especially true if we combine it with Psyco.

14.1 Exercise

Run the comparison with different values for `total` and `repeat`. Use very small and very large numbers. Compare the ratios between the versions. Explain the differences. Use PyPy for the tests. Change the number of used processes in the multiprocessing examples.

15 Compilation of Tools for Speedup and Extending

There are many more ways to extend Python. Therefore, a short compilation of methods and tools for this purpose is given here. The compilation is by no means exhaustive.

Method/Tool	Remarks	Link
algorithmic improvements	try this first	http://www.python.org
Numpy	matlab like array processing	http://numpy.scipy.org
Psyco	JIT, very little to do	http://psyco.sourceforge.net
PyPy	fast Python implementation	http://pypy.org/
Cython	C with Python syntax	http://cython.org/
ctypes	call DLLs directly	Python's standard library
C extensions by hand	lots of work	http://docs.python.org/ext/ext.html
SWIG	mature, stable, widely used	http://www.swig.org
Boost.Python	C ++ template based, elegant	http://www.boost.org/libs/python/doc
PyCXX/CXX	last update August 2005	http://cxx.sourceforge.net
SCXX	last update 2002	http://davidf.sjsoft.com/mirrors/mcmillan-inc/scxx.html
SIP	developed for Qt, fast	http://www.riverbankcomputing.co.uk/sip
Babel	unite C/C++, F77/90, Py, Java	http://www.llnl.gov/CASC/components
Weave	inline C++ in Python	http://scipy.org/Weave
PyInline	inline other languages (alpha)	http://pyinline.sourceforge.net
f2py	stable Fortran extensions	http://cens.ioc.ee/projects/f2py2e
pyFort	F77 only, last update 2004	http://pyfortran.sourceforge.net
PyCUDA	Python on GPGPUs	http://mathemat.tician.de/software/pycuda
PyOpenCL	Python on GPGPUs	http://mathemat.tician.de/software/pyopencl
Copperhead	Python on GPGPUs	http://code.google.com/p/copperhead/
COM/DCOM, CORBA, XML-RPC, ILU	middleware	various

16 End

16.1 Colophon

This material is written in `reStructuredText` and has been converted into PDF using `rst2pdf`. All Python files are dynamically included from the development directory.

16.2 Links

LANG2006	Hans Petter Lantangen, Python Scripting for Computational Science, Second Edition, Springer Berlin, Heidelberg, 2006.
MART2005	Alex Martelli et al., Python Cookbook, O'Reilly, 2nd Edition, 2005.
MART2006	Alex Martelli, Python in a Nutshell, O'Reilly, 2nd Edition, 2006.
ZIAD2008	Tarek Ziadè, Expert Python Programming: Best practices for designing, coding, and distributing your Python software, Packt 2008.
1	http://aspn.activestate.com/ASPN/Python/Cookbook
2	http://www.vrplumber.com/programming/runsnakerun/
3	http://guppy-pe.sourceforge.net
4	http://packages.python.org/Pympler/
5	http://en.wikipedia.org/wiki/Big_O_notation
6	http://www.danga.com/memcached/
7	http://docs.djangoproject.com/en/dev/topics/cache/
8	http://psyco.sourceforge.net
9	http://numpy.scipy.org